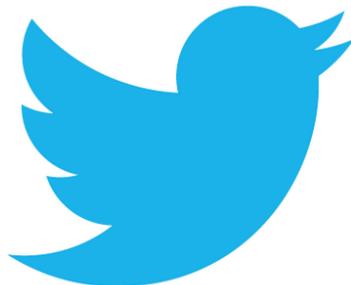


Parallel Programming Principle and Practice

Lecture 9 – Programming for Big Data Processing

What is big data?

A massive volume of both structured and unstructured data that is too large to process with traditional database and software techniques.



- 34K likes every minute
- It deals with 3-4 PB data every day
- 1 billion active user
- It generates 12TB data daily
- 200 million users generate 230 million tweets daily
- 2 million search every minute
- It deals with 20 PB data every day

Application of big data analytics

Smart healthcare



Manufacturing



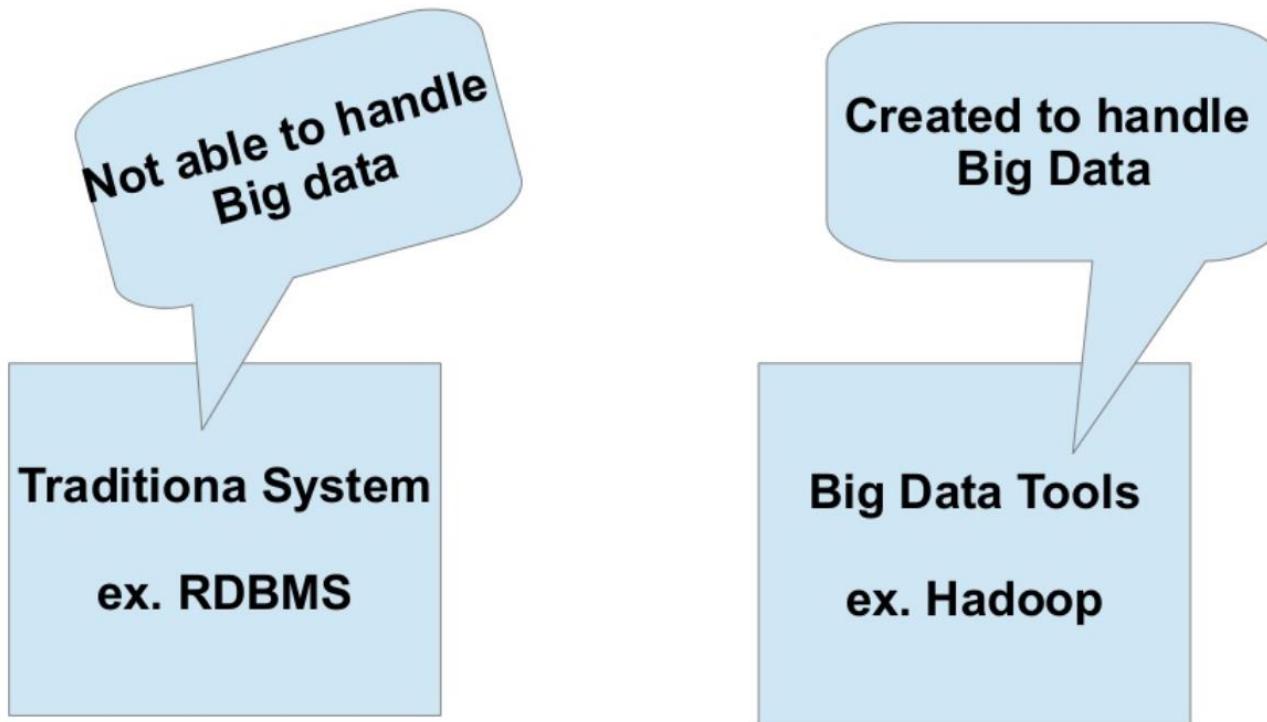
Traffic control



Trading analytics



Tools for handling big data analytics



Outline

- Programming for batch data processing
- Programming for graph processing
- Programming for stream processing

PROGRAMMING FOR BATCH DATA PROCESSING

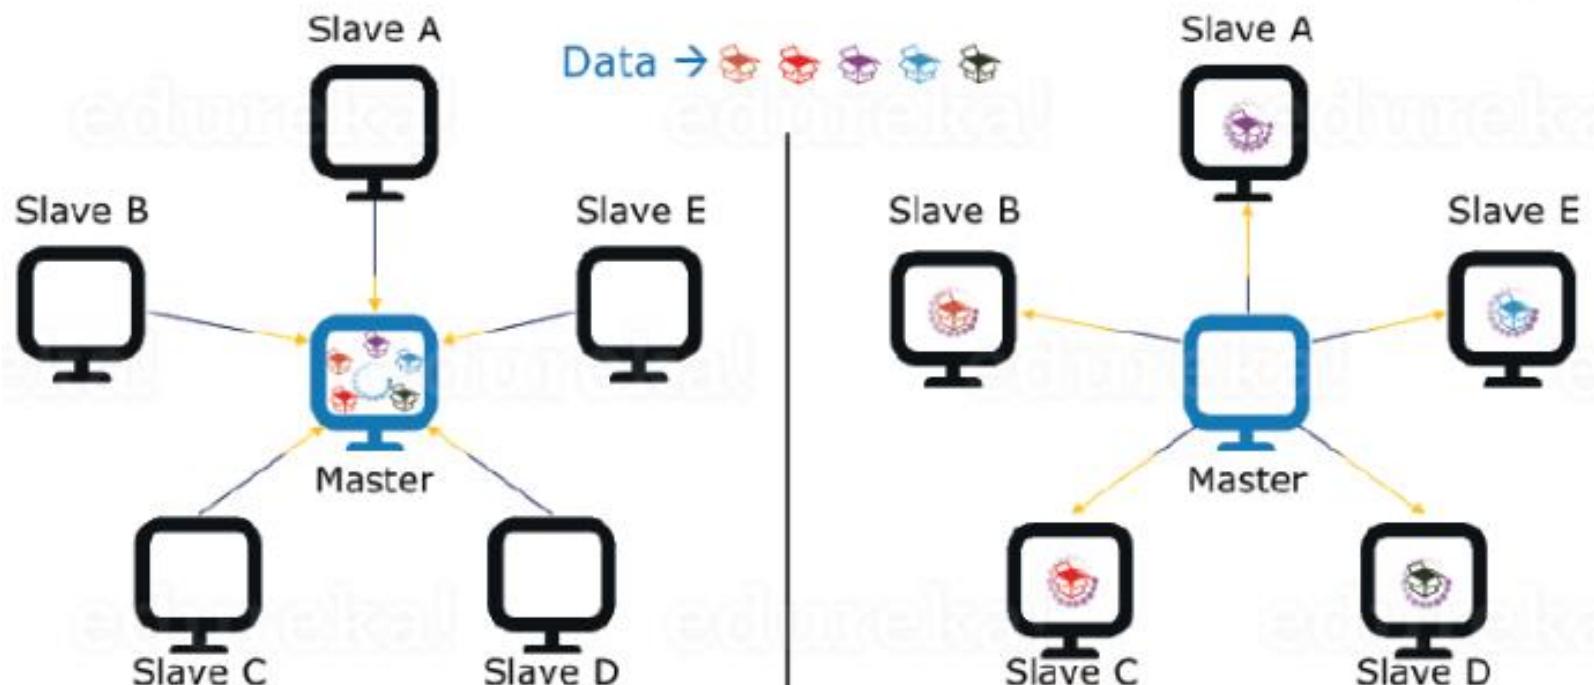
Outline

-
- Background of batch data processing
 - Programming model for batch data processing
 - MapReduce framework

Batch data processing

- Big data analytics usually contain tasks that perform batch data processing
 - Tasks that can be executed offline without user interaction
 - It process all data records in batch
 - Tools for batch processing of big data
- Challenge
 - Data placement, computation overhead
- Solution
 - Parallel execution on clusters with thousands of machines
 - Moving processing unit to the data

Big data placement

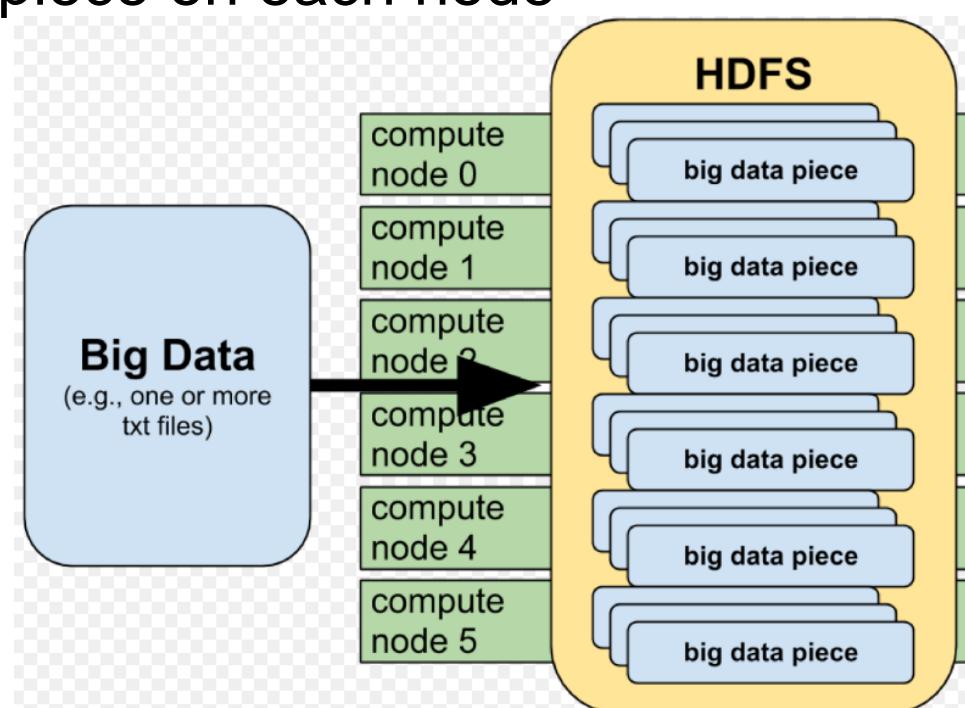


1. Moving data to the Processing Unit
(Traditional Approach)

2. Moving Processing Unit to the data
(MapReduce Approach)

Parallel batch data processing

- Divide the input data into pieces
- Distribute the pieces on nodes and manage the data with distributed file system
- Parallel process of data piece on each node
- Aggregate the results



Challenges for programmers

- Data placement
 - Data division, redundancy
- Computing parallelization
 - Synchronization between processes, consistency
- Communication
- Fault tolerant
- Load balancing

Solution

- A general abstraction for special-purpose applications
 - Express the various simple computation
 - Hide the details of parallelization, such as fault-tolerance, data distribution, and load balancing
 - Provide powerful interfaces to users
 - Programmers can implement the application code without considering the messy parallelization details
 - Automatic parallelization and distribution of large-scale computation

A general programming interface for various special-purpose applications is required!

Outline

-
- Background of batch data processing
 - Programming model for batch data processing
 - MapReduce framework

Programming model for batch data processing

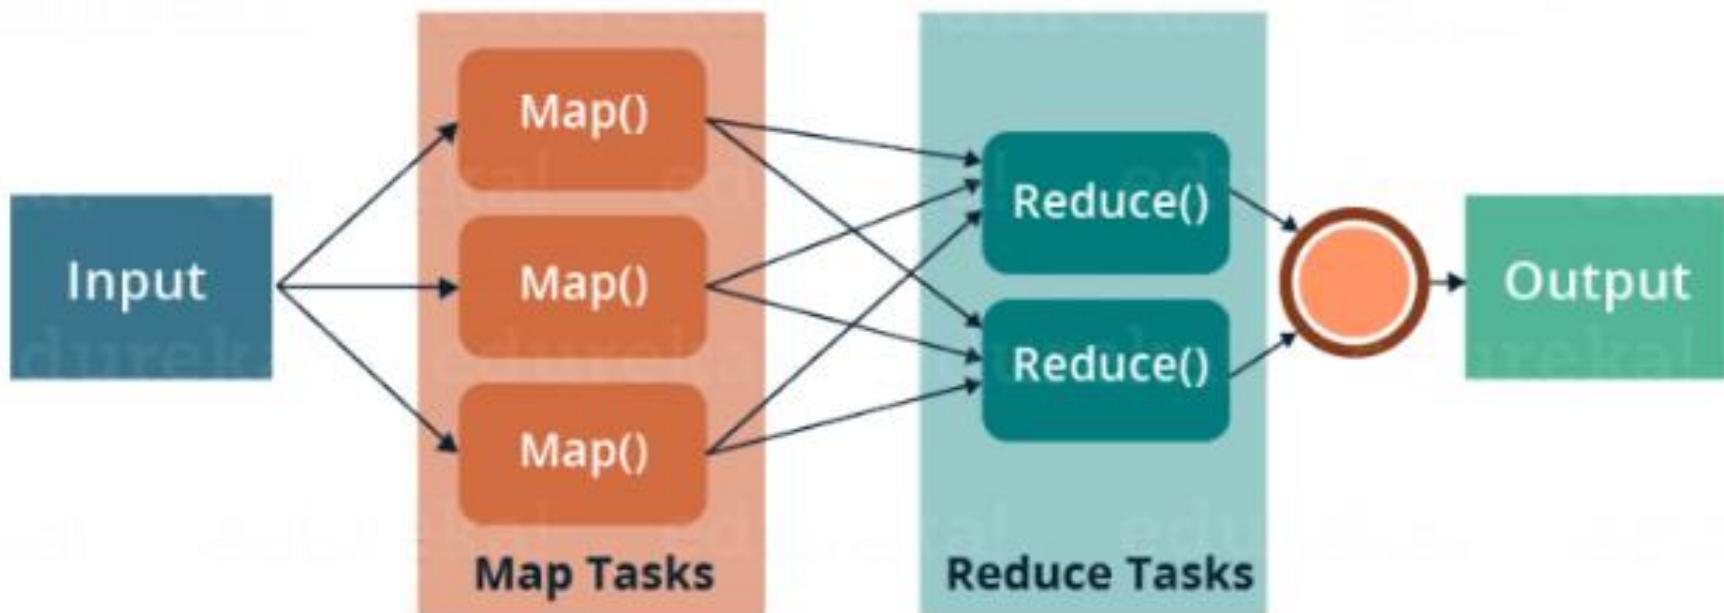
- Inspired by primitives of Lisp and other functional languages
- Most of our computations can be represented by a map and reduce operator to each logical record in the input dataset
 - The map and reduce operator are interfaces for programmers
- We can develop a general framework **MapReduce** for batch data processing
 - Provide the powerful map/reduce interface for programmers
 - Deal with parallelization details transparently and automatically

Program model

- Read a lot of data
- **Map**: extract something you care about from each data record
- Shuffle and sort the data
- **Reduce**: aggregate, summarize, filter, or transform the intermediate data
- Write the final results

Users can change the definition of map and reduce to fit different problems

Program model



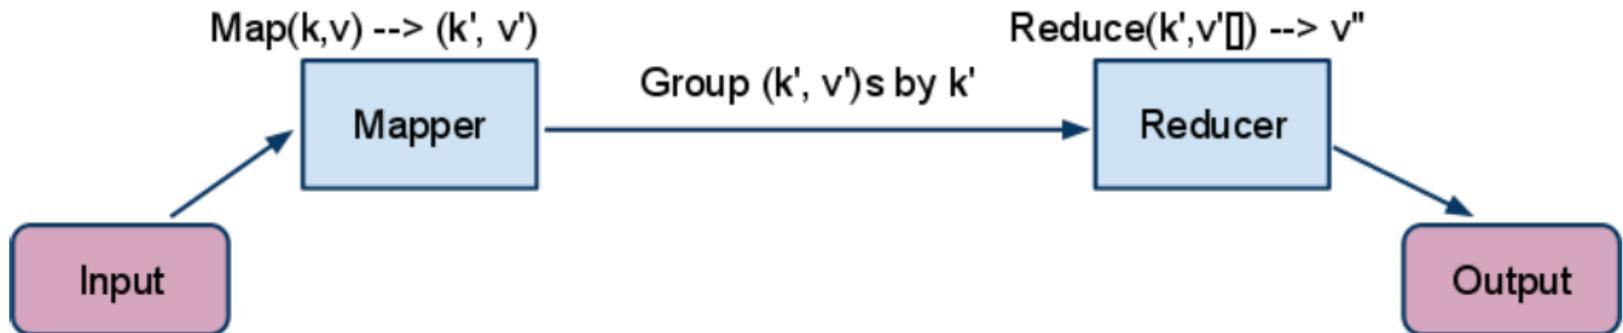
Map/reduce abstraction for large-scale data

- Divide the large-scale input dataset into splits
- Distribute all splits on machines
- Perform map operator on each split
- Shuffle and sort intermediate results
- Perform reduce operator on each sorted results
 - Aggregate the intermediate results
- Generate final output

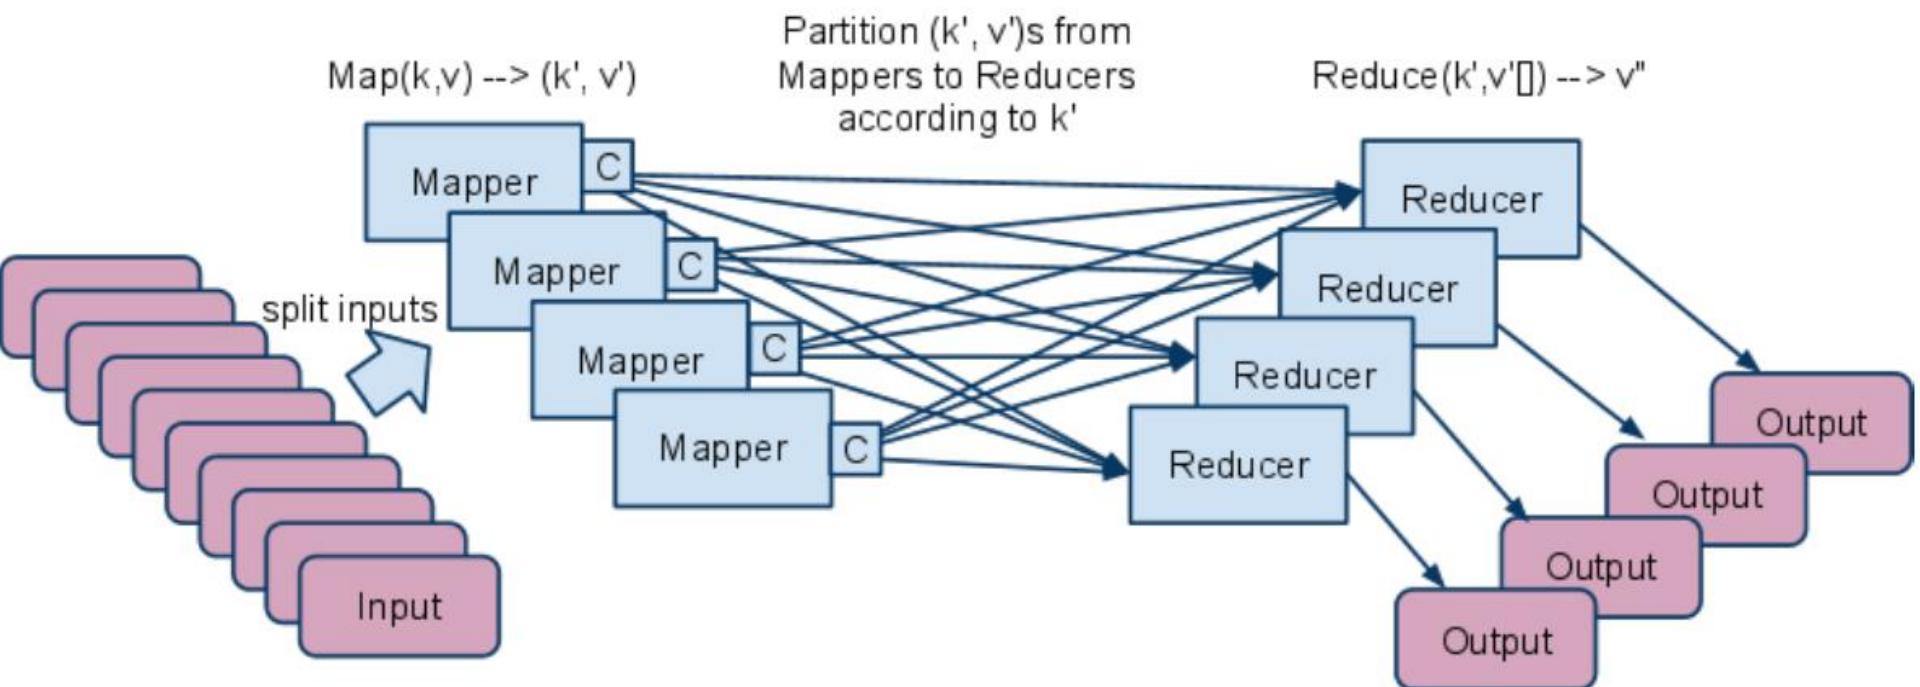
Programmers only need to specify the map and reduce operator without considering other details.

Map/reduce operators

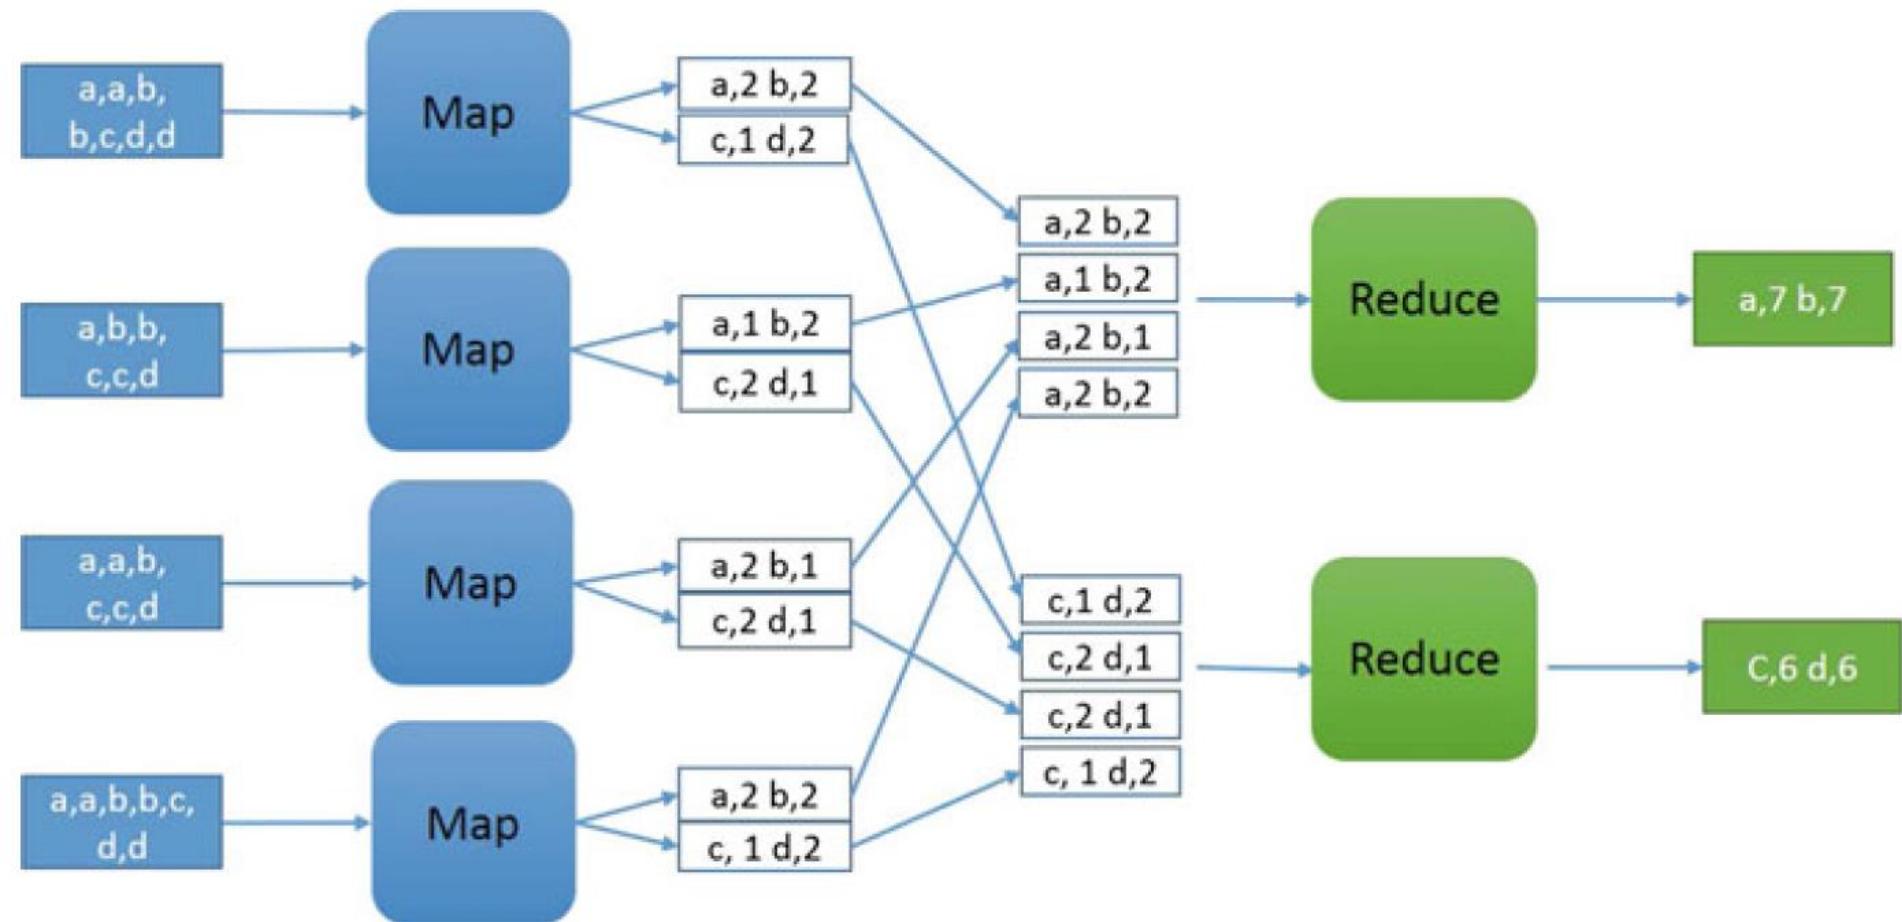
- Programmers specify two primary functions
 - $\text{map } (k, v) \rightarrow \langle k', v' \rangle$
 - $\text{reduce } (k', [v']) \rightarrow \langle k'', v'' \rangle$
 - All v' with the same k' are sent to the same reducer (shuffle)



Map/reduce abstraction for large-scale data



Word count solution



Word count solution

```
//Pseudo-code for "word counting"
map(String key, String value) :
    // key: document name,
    // value: document contents
    for each word w in value:
        EmitIntermediate(w, "1");

reduce(String key, Iterator values) :
    // key: a word
    // values: a list of counts
    int word_count = 0;
    for each v in values:
        word_count += ParseInt(v);
    Emit(key, AsString(word_count));
```

Outline

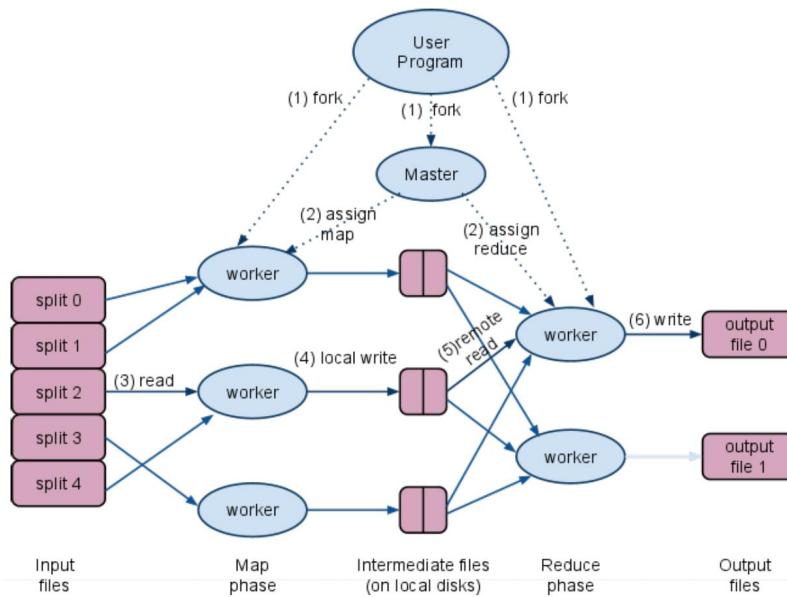
- Batch data processing
- Programming model for batch data processing
- MapReduce framework

Functions of MapReduce framework

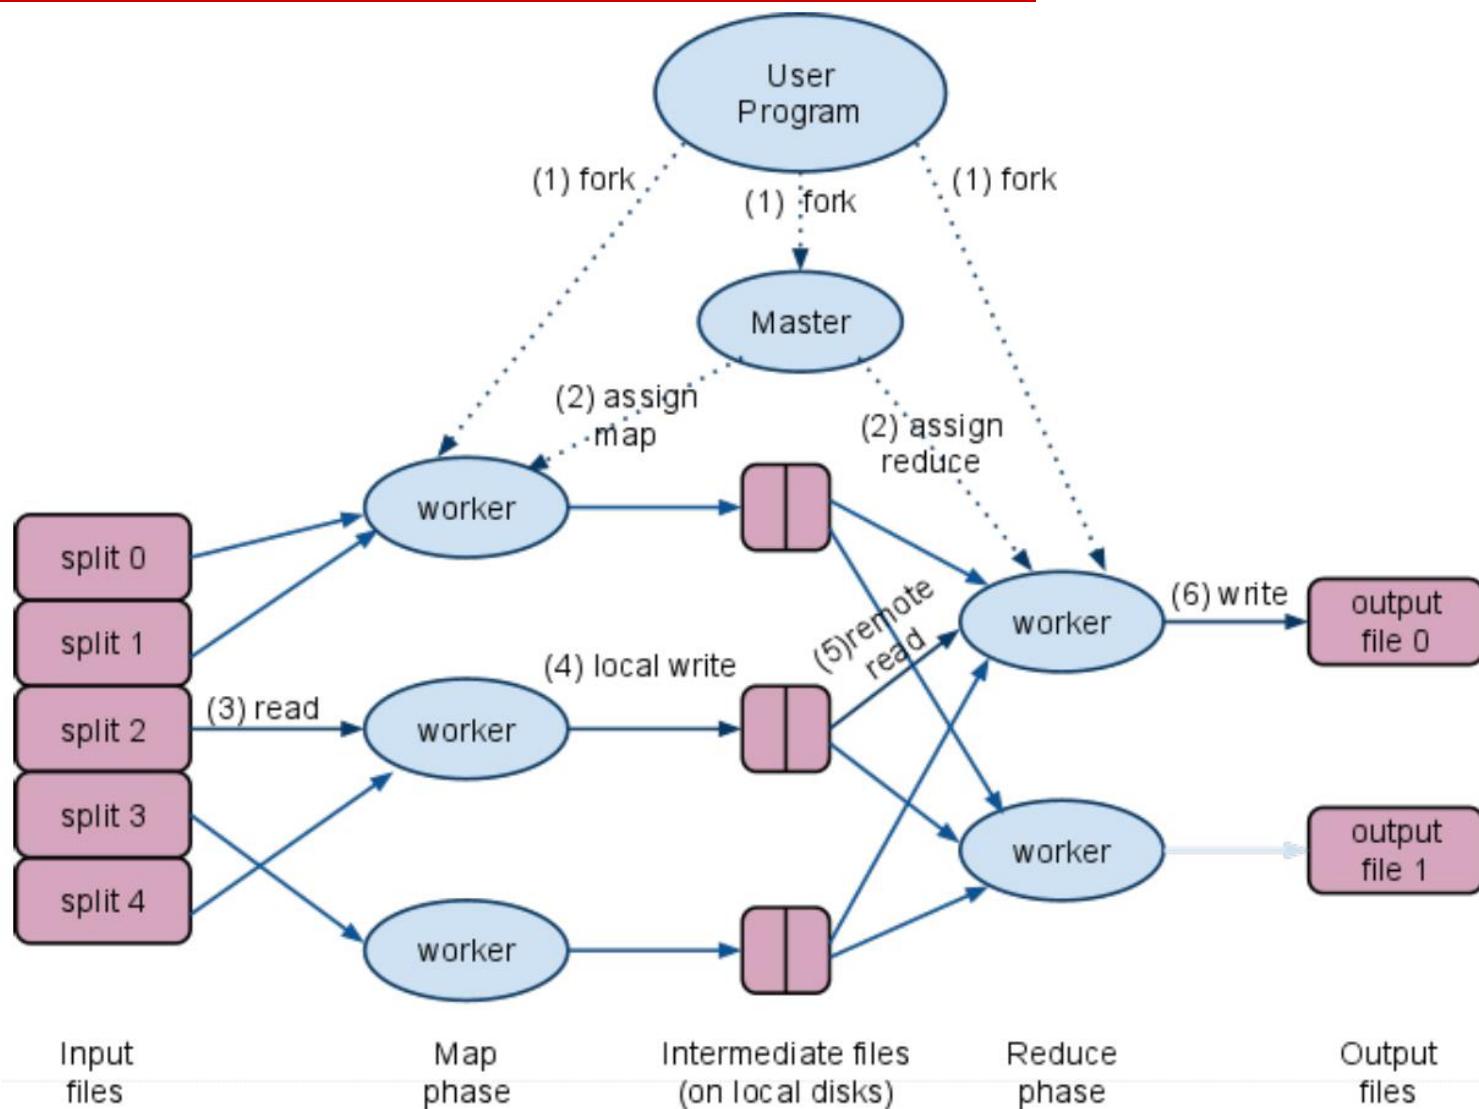
- Programmers specify the map/reduce operators, MapReduce framework should handle
 - Data distribution
 - Split data for multiple workers, move processes to data
 - Synchronization
 - Shuffle intermediate data
 - Fault tolerance
 - Detect failed machines and automatically restart, recover the failed tasks
 - Scheduling
 - Load balancing, detect stragglers
 - Scalability
 - Fast growing dataset, extension of machines

Architecture of MapReduce framework

- One master, many workers
- Master assigns map/reduce task to a free worker
- Workers perform the assigned tasks with corresponding input data
- Master schedules the tasks, handles the failed workers



Execution overview of MapReduce framework



Step 1: split input files into shards

- Break up the input data into M pieces (typically 64MB)

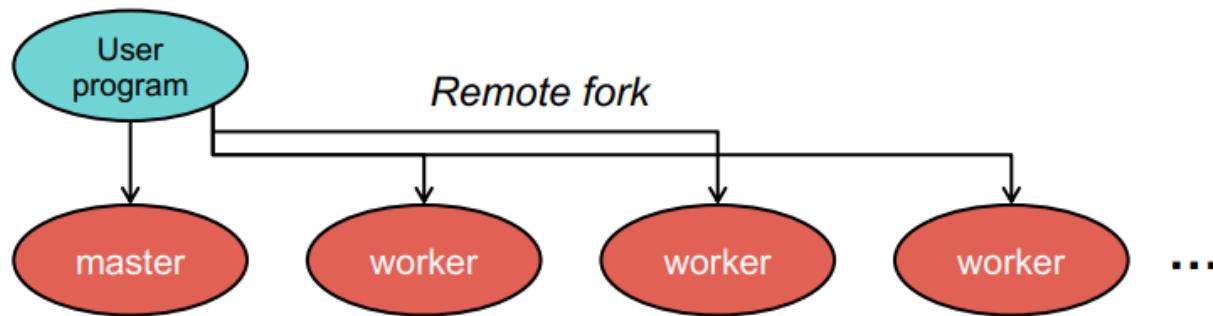


Input files

Divided into M shards

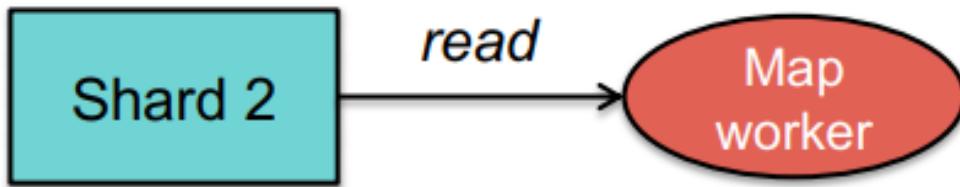
Step 2: fork processes

- Start up many copies of the program on a cluster of machines
 - One master: scheduler and coordinator
 - Lots of workers
- Idle workers are assigned
 - Map tasks (each works on a shard) – there are M map tasks
 - Reduce tasks (each works on intermediate pairs) – there are R reduce tasks
 - R is defined by the user



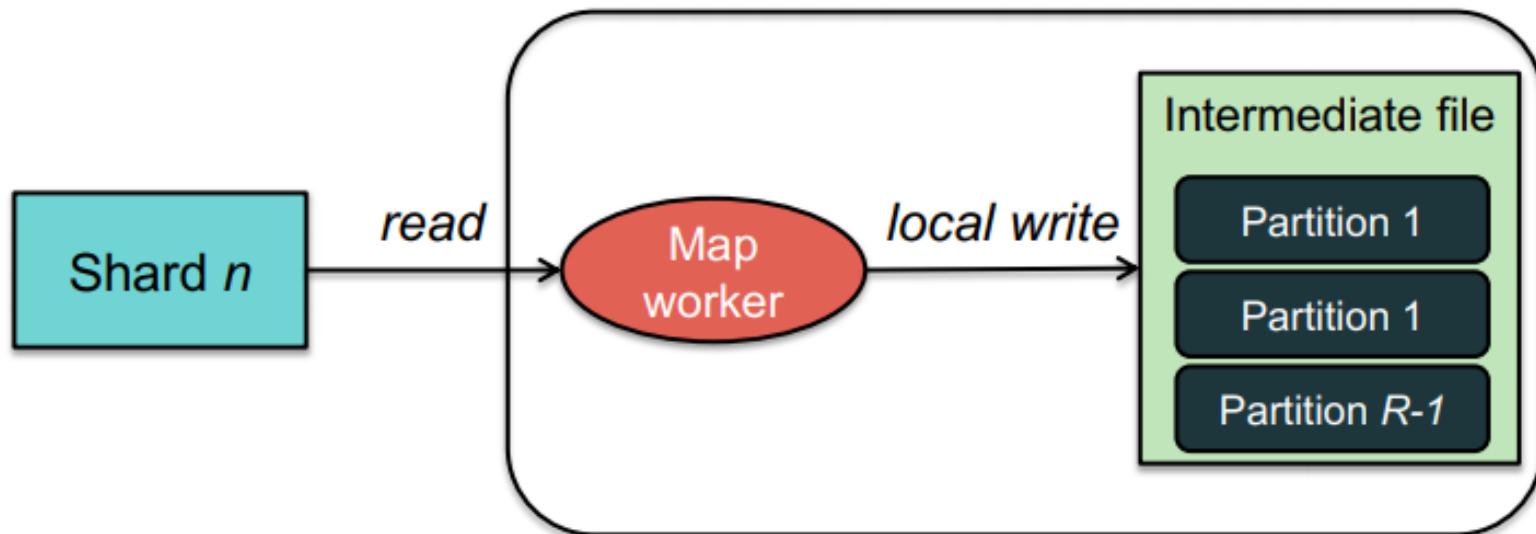
Step 3: run map tasks

- Map worker reads the input shard assigned to it
- Parse key/value pairs out of the input shard
- Pass each pair to the user-specified map function
 - Produce intermediate key/value pairs
 - Intermediate pairs are buffered in memory



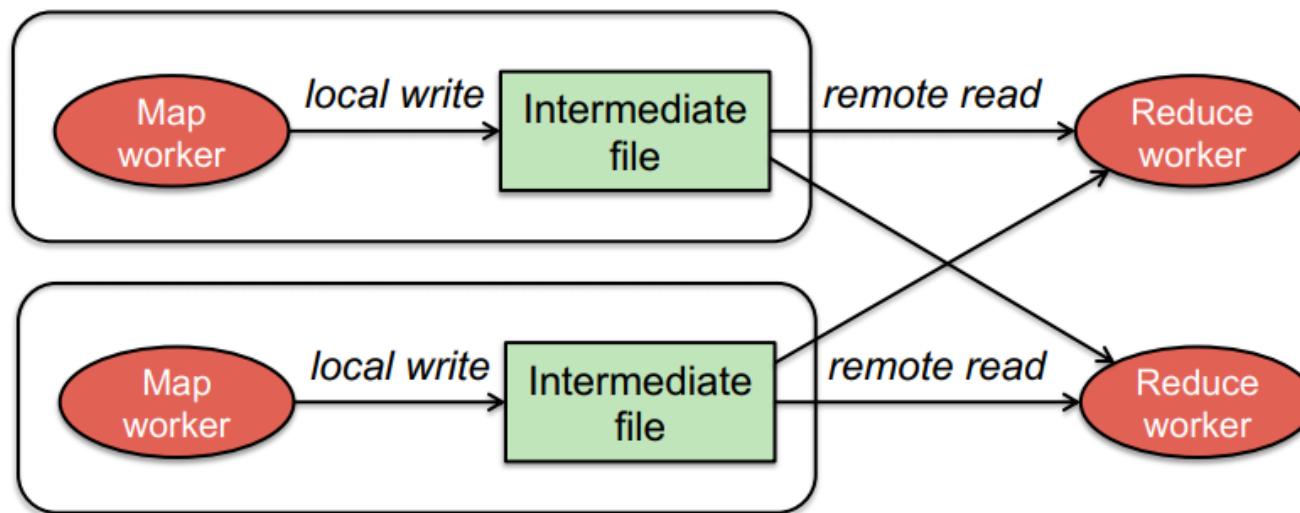
Step 4: create intermediate files

- Intermediate pairs produced by map worker are periodically written to the local disk
 - Partitioned into R regions by a partitioning function
 - Sort all intermediate pairs by keys



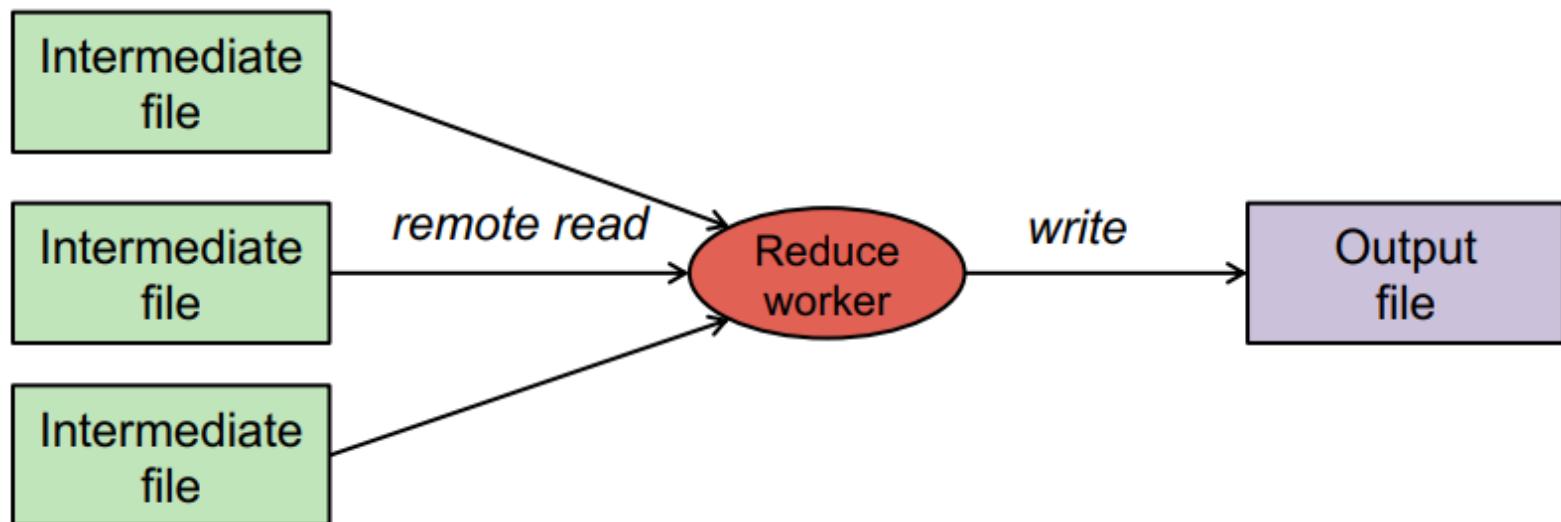
Step 5: sorting

- Reduce worker gets notified by the master about the location of intermediate files for its partition
- Read the data from the local disk of the map workers
- Sorts the data by the intermediate keys
- All pairs with the same key are grouped together



Step 6: run reduce tasks

- Pass the sorted and grouped pairs to the user-specified reduce function
- The output of reduce function is appended to an output file



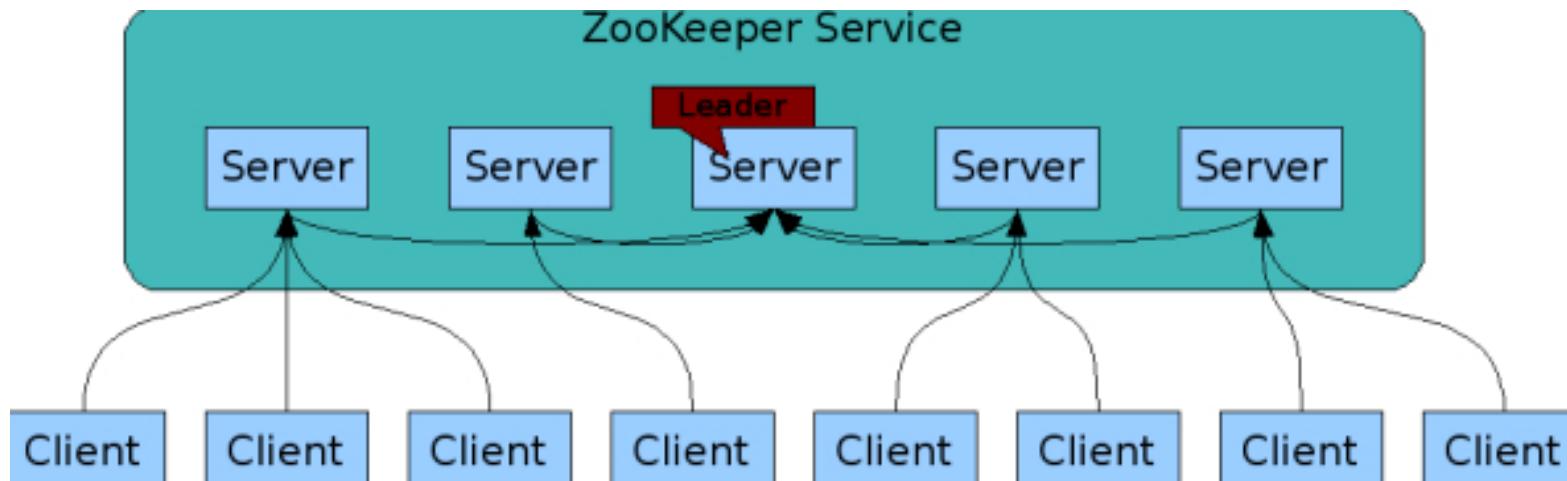
Step 7: return to user

- When all map and reduce tasks have completed, the master wakes up the user program
 - The output of MapReduce is saved in R files
- The MapReduce call in the user program returns

Coordinate

- A centralized service for MapReduce framework

- Maintain configuration information of the framework
- Provide distributed synchronization between processes
- Guarantee the consistency of the framework
- Apache Zookeeper, Google Chubby, etc.



Fault tolerance

□ Worker failure

- Detect failure via periodic heartbeats
- Master reschedules the tasks to other workers
- Re-execute completed and in-progress map tasks
- Re-execute in-progress reduce tasks
- Task completion committed through master

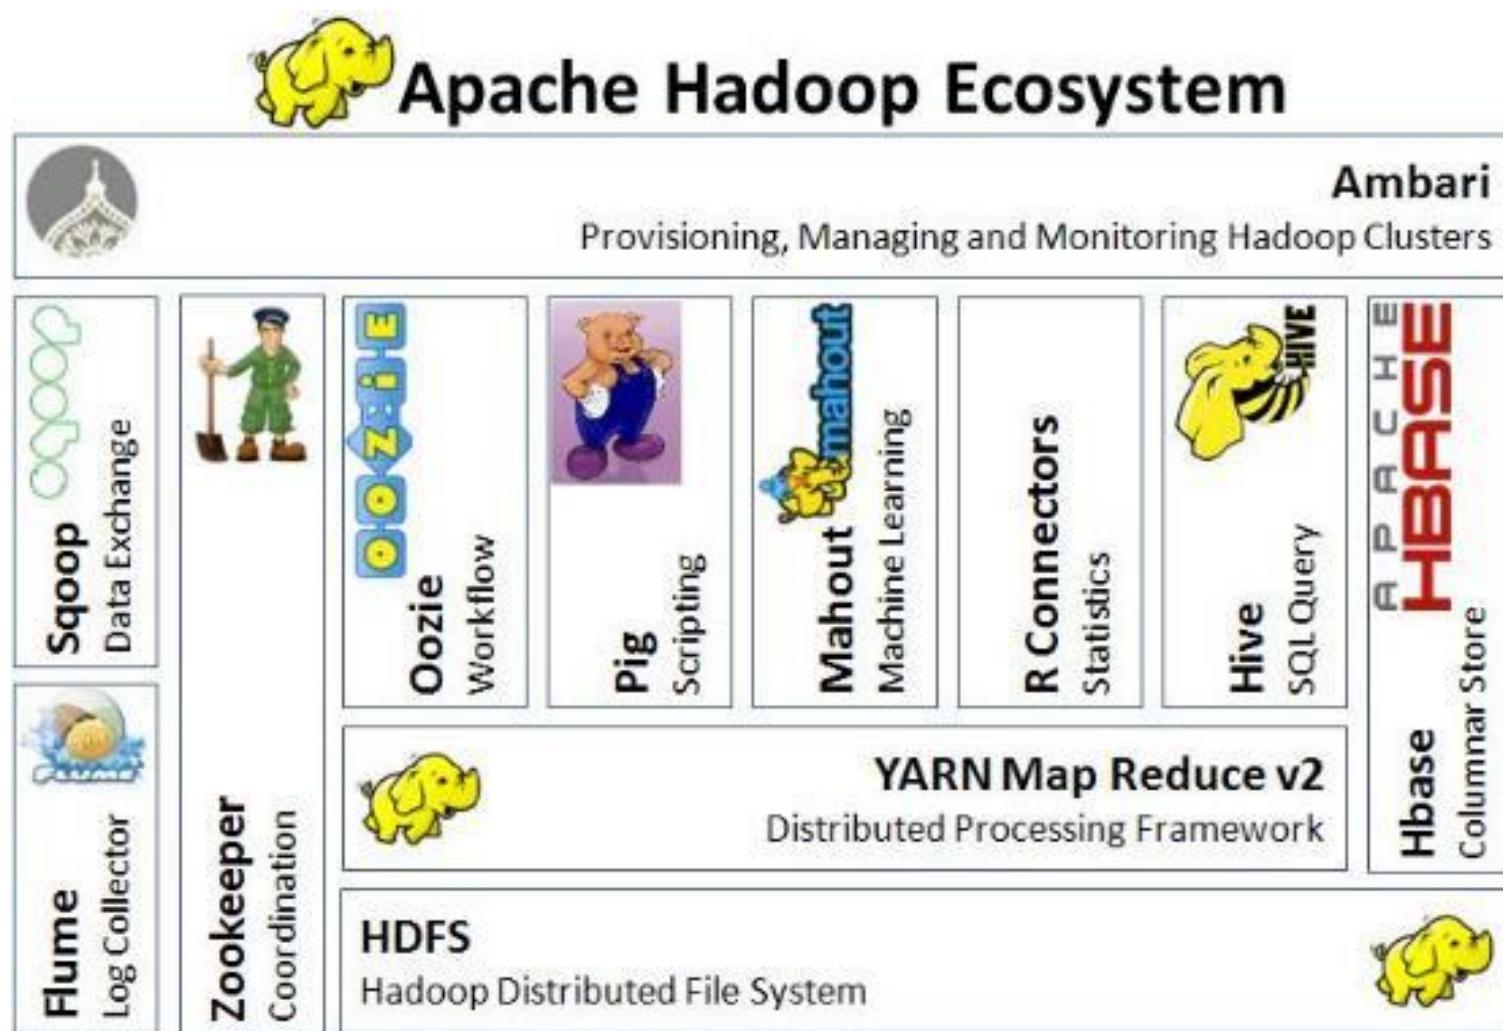
□ Master failure

- Select a new master
- Recover the master according to the checkpoint on distribute filesystem

MapReduce implementations

- Google has a proprietary implementation in C++
- Hadoop is an open-source implementation in Java
 - Developed by Yahoo, used in production
 - An Apache project
 - Rapidly expanding software ecosystem
- Lots of custom research implementations
 - For GPUs, FPGAs, etc.

Ecosystem of Hadoop



PROGRAMMING FOR GRAPH PROCESSING

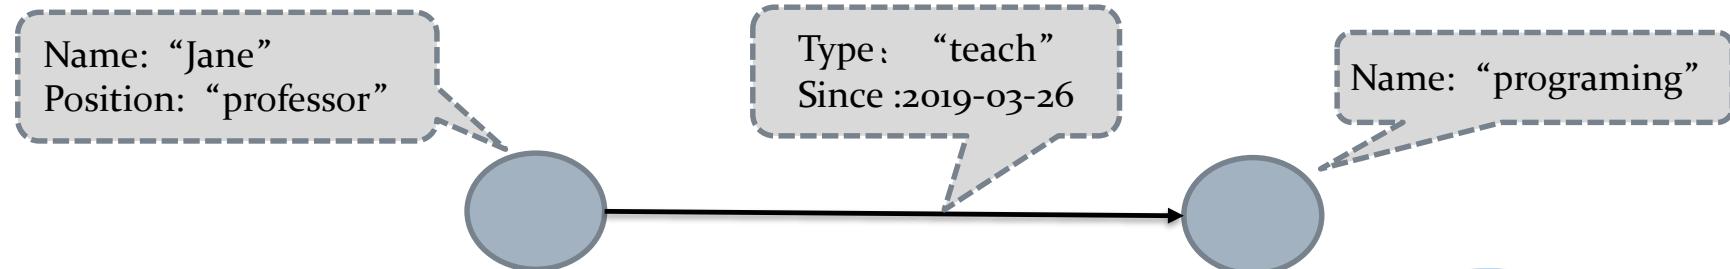
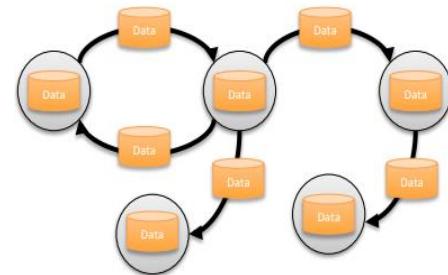
What is graph?

- A graph is a collection of **vertices** connected by **edges** ($G = (V, E)$)

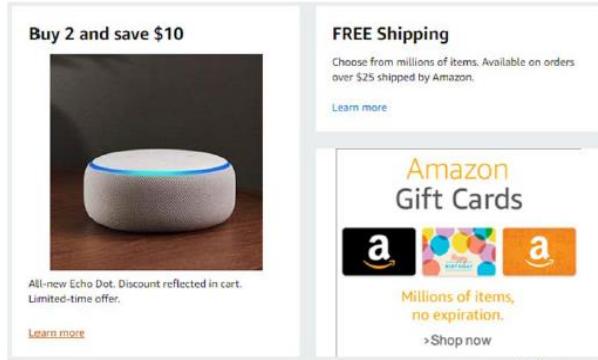
- Edge may be **directed** (from A to B)

or **undirected** (between A and B)

- Vertices and edges may have **properties**



Graph processing is universal



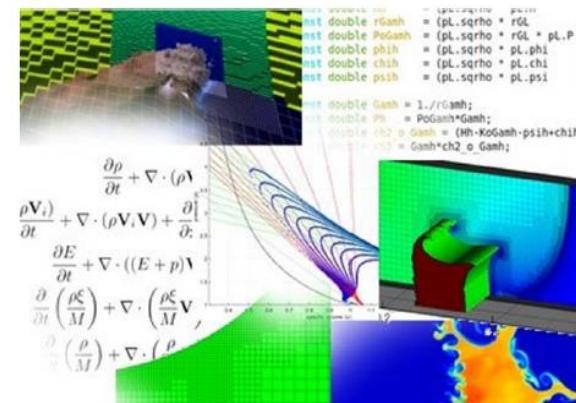
advertising recommendation



social network analysis



network ranking



scientific calculation

Graph size is increasing



NETFLIX

> 1 M vertices
> 100 M edges
*Distributed
graphLab-2012

advertising network



> 1 B vertices
> 1 T edges
*Facebook
Engineering Blog

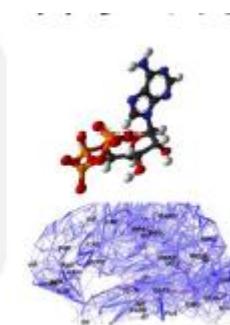
social media



> 50 B vertices
> 1 T edges
*NSA Big Graph
Experiment -
2013

internet

> 100 B vertices
> 100 T edges
*NSA Big graph
Experiment -
2013



scientific calculation

Challenges of graph processing

Data-driven computation

- Computation for graph task is closely related to vertex and edge
- Computation structure is hard to predict before running, hard to parallel

Irregular structure

- Hard to divide the irregular structure of graph with good quality

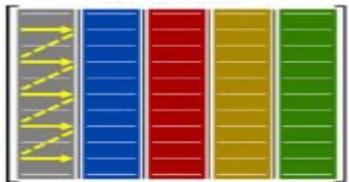
Poor access locality

- Non-contiguous access for vertex/edge
- Existing caching mechanisms can only accelerate access with good locality

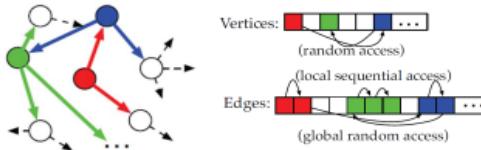
Access/computation ratio is high

- Large number of access lead to CPU waiting

Graph processing vs traditional processing



- **Compute intensive**
Computation complexity is high, easy to cover memory latency
- **Sequential memory access**
Data is stored in memory sequentially
- **High data parallelism**
No complex dependency between different data, convenient for parallel processing
- **Good locality**
With good spatial and time locality

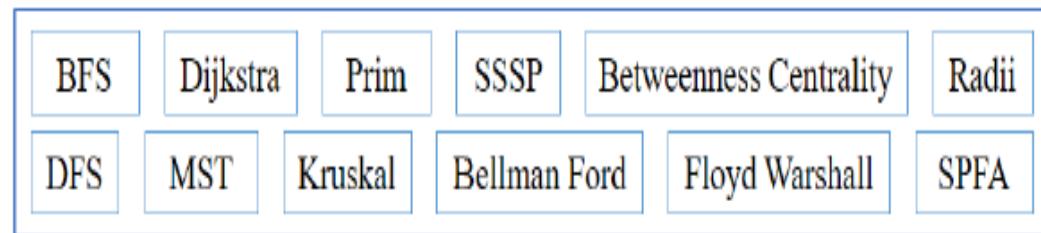


- **Low computation/access ratio**
Little computation in each node, hard to cover latency
- **Large random access**
There is a large number of random access requests across the region
- **Complex data dependency**
Difficult to explore parallelism, a lot of data conflicts
- **Unstructured distribution**
Uneven distribution of vertex degree, serious load balancing problems and communication overhead

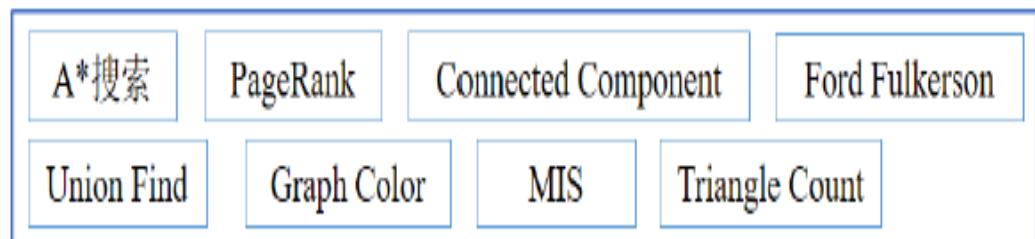
Graph algorithm

- **Traversal-centric algorithms**
 - Require a specific way to traverse the graph from a particular vertex
 - Have a large number of random access
- **Computation-centric algorithms**
 - A large number of operations in each iteration
 - All nodes participate in each iteration

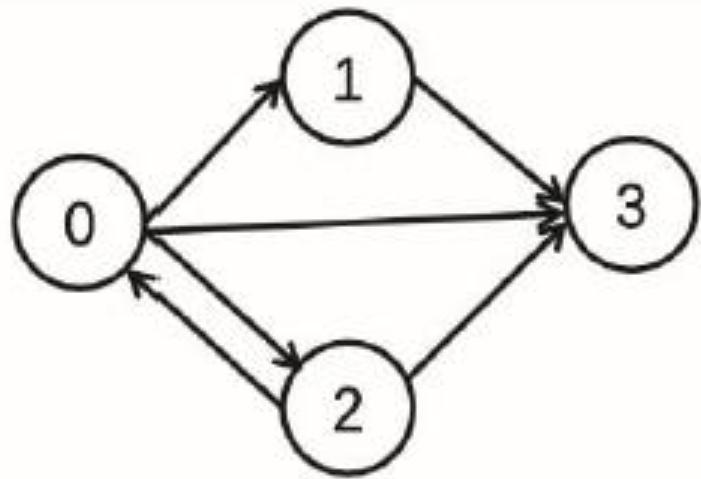
Traversal-centric algorithms



Computation-centric algorithms

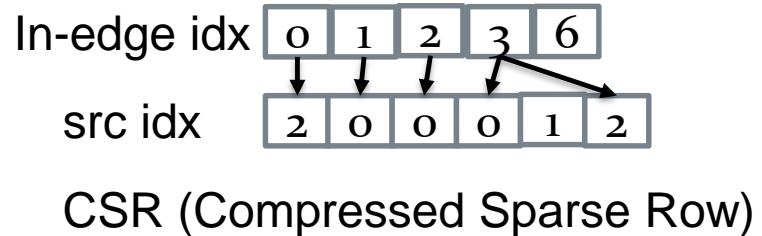
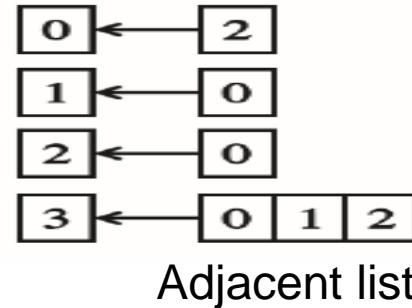
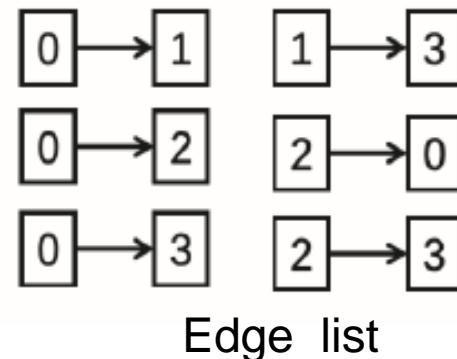


Graph storage



Graph Topology

Access efficiency vs Compact ratio



Graph processing frameworks

Programming Model

Vertex Centric: Pregel^[SIGMOD'10]
Edge Centric: X-Stream^[SOSP'13]
Path Centric: PathGraph^[SC'14]
Subgraph centric: Blogel^[VLDB'14]

Graph Partition Strategy

Vertex Cut: Pregel^[SIGMOD '10]
Edge cut: PowerGraph^[OSDI'12]
Hybrid Cut: PowerLyra^[EuroSys '15]

Computation Strategy

Pull vs push
Synchronous vs Asynchronous

Typical System

Pregel : SIGMOD 2010
PowerGraph: OSDI 2012
GraphChi: OSDI 2012

Graph programming model

Programming model

- Easy and concise to use
- Flexible to express different graph tasks
- Efficient to fit in computation architecture

Existing programming model

- Vertex-centric
- Edge-centric
- Path-centric
- Subgraph-centric

Programming model: vertex-centric

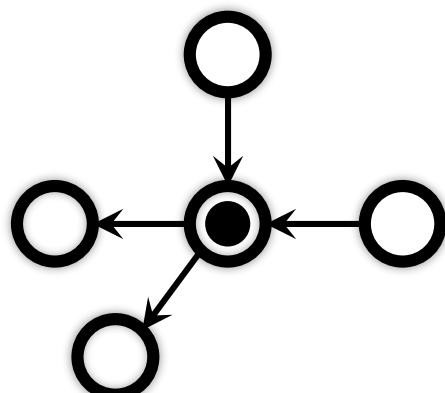
Coding graph algorithms as **vertex-centric** programs to process **vertices** in parallel and communicate along **edges**

-- "Think as a Vertex" philosophy

Example: *PageRank*

A centrality analysis algorithm to iteratively rank each vertex as a weighted sum of neighbors' ranks

$$R_i = 0.15 + 0.85 \sum_{(j, i) \in E} \omega_{ij} R_j$$



```

COMPUTE(v)
foreach n in v.in_nbrs
    sum += n.rank / n.nedges
v.rank = 0.15 + 0.85 * sum
if !converged(v)
    foreach n in v.out_nbrs
        activate(n)
    
```

Programming model: vertex-centric

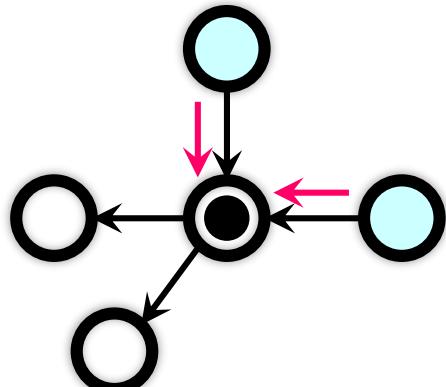
Coding graph algorithms as **vertex-centric** programs to process **vertices** in parallel and communicate along **edges**

-- "Think as a Vertex" philosophy

Example: *PageRank*

A centrality analysis algorithm to **iteratively** rank each vertex as a weighted sum of neighbors' ranks

$$R_i = 0.15 + 0.85 \sum_{(j, i) \in E} \omega_{ij} R_j$$



Gather data from neighbors via in-edges

COMPUTE(v)

```
foreach n in v.in_nbrs
    sum += n.rank / n.nedges
```

v.rank = 0.15 + 0.85 * sum

```
if !converged(v)
    foreach n in v.out_nbrs
        activate(n)
```

Programming model: vertex-centric

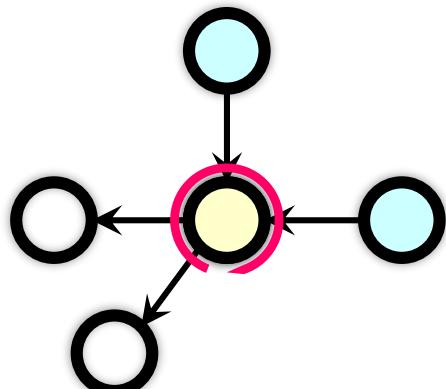
Coding graph algorithms as **vertex-centric** programs to process **vertices** in parallel and communicate along **edges**

-- "Think as a Vertex" philosophy

Example: *PageRank*

A centrality analysis algorithm to **iteratively** rank each vertex as a weighted sum of neighbors' ranks

$$R_i = 0.15 + 0.85 \sum_{(j, i) \in E} \omega_{ij} R_j$$



Update vertex with
new data

```

COMPUTE(v)
foreach n in v.in_nbrs
    sum += n.rank / n.nedges

v.rank = 0.15 + 0.85 * sum

if !converged(v)
    foreach n in v.out_nbrs
        activate(n)
    
```

Programming model: vertex-centric

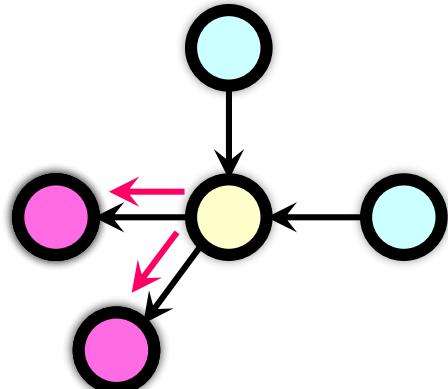
Coding graph algorithms as **vertex-centric** programs to process **vertices** in parallel and communicate along **edges**

-- "Think as a Vertex" philosophy

Example: *PageRank*

A centrality analysis algorithm to **iteratively** rank each vertex as a weighted sum of neighbors' ranks

$$R_i = 0.15 + 0.85 \sum_{(j, i) \in E} \omega_{ij} R_j$$



Scatter data to neighbors via out-edges

```

COMPUTE(v)
foreach n in v.in_nbrs
    sum += n.rank / n.nedges

v.rank = 0.15 + 0.85 * sum

if !converged(v)
    foreach n in v.out_nbrs
        activate(n)
    
```

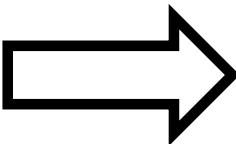
Programming model: edge-centric

Coding graph algorithms as **edge-centric** programs to process **edges** in parallel, where **edge** is the minimum parallel processing unit.

-- "Think as an Edge" philosophy

```
for each vertex v
  if v has update
    for each edge e from v
      scatter update along e
```

Scatter



```
for each edge e
  If e.src has update
    scatter update along e
```

Scatter

Vertex-Centric

Transformation

Edge-Centric

- Balanced load for fine-grained parallel processing unit
- Better locality for sequential access of the edges

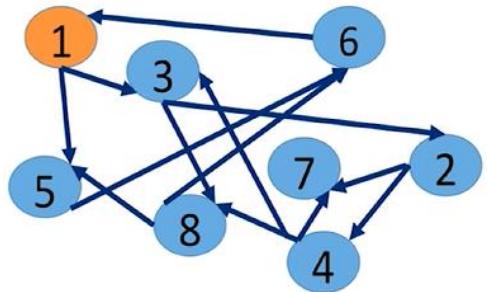
Programming model: edge-centric

Coding graph algorithms as **edge-centric** programs to process **edges** in parallel, where **edge** is the minimum parallel processing unit.

-- "Think as an Edge" philosophy

Example: BFS

Edge-Centric Scatter Gather



BFS



V
1
2
3
4
5
6
7
8

Edge-Centric Scatter Gather for BFS

SOURCE	DEST
1	3
1	5
2	7
2	4
3	2
3	8
4	3
4	7
4	8
5	6
6	1
8	5
8	6

Edge-centric Scatter-Gather

edge_scatter(edge e)
send update over e

update_gather(update u)
apply update u to u.destination

while not done
for all edges e
 edge_scatter(e)
for all updates u
 update_gather(u)

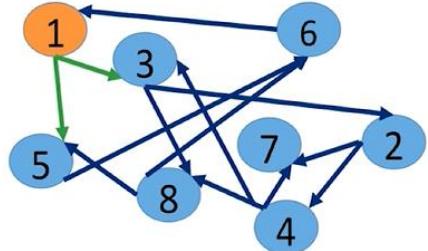
Programming model: edge-centric

Coding graph algorithms as **edge-centric** programs to process **edges** in parallel, where **edge** is the minimum parallel processing unit.

-- "Think as an Edge" philosophy

Example: BFS

Edge-Centric Scatter Gather



BFS



v
1
2
3
4
5
6
7
8

SOURCE	DEST
1	3
1	5
2	7
2	4
3	2
3	8
4	3
4	7
4	8
5	6
6	1
8	5
8	6

Edge-Centric Scatter Gather for BFS

Edge-centric Scatter-Gather

edge_scatter(edge e)
send update over e

update_gather(update u)
apply update u to u.destination

while not done
for all edges e
 edge_scatter(e)
 for all updates u
 update_gather(u)

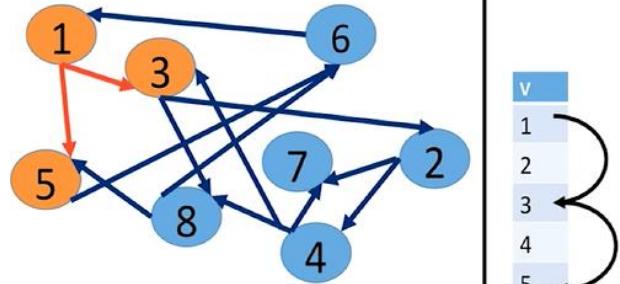
Programming model: edge-centric

Coding graph algorithms as **edge-centric** programs to process **edges** in parallel, where **edge** is the minimum parallel processing unit.

-- "Think as an Edge" philosophy

Example: BFS

Edge-Centric Scatter Gather



BFS

Edge-Centric Scatter Gather for BFS



SOURCE	DEST
1	3
1	5
2	7
2	4
3	2
3	8
4	3
4	7
4	8
5	6
6	1
7	5
8	6

Edge-centric Scatter-Gather

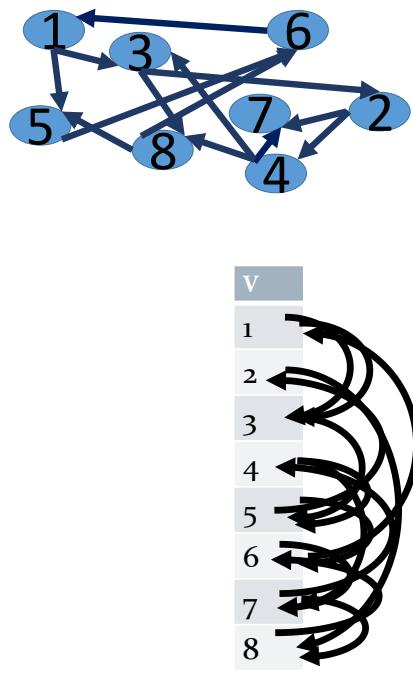
edge_scatter(edge e)
send update over e

update_gather(update u)
apply update u to u.destination

```
while not done
  for all edges e
    edge_scatter(e)
  for all updates u
    update_gather(u)
```

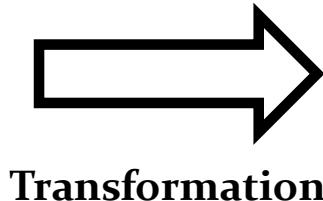
Programming model: edge-centric

Coding graph algorithms as **edge-centric** programs to process **edges** in parallel, where **edge** is the minimum parallel processing unit.

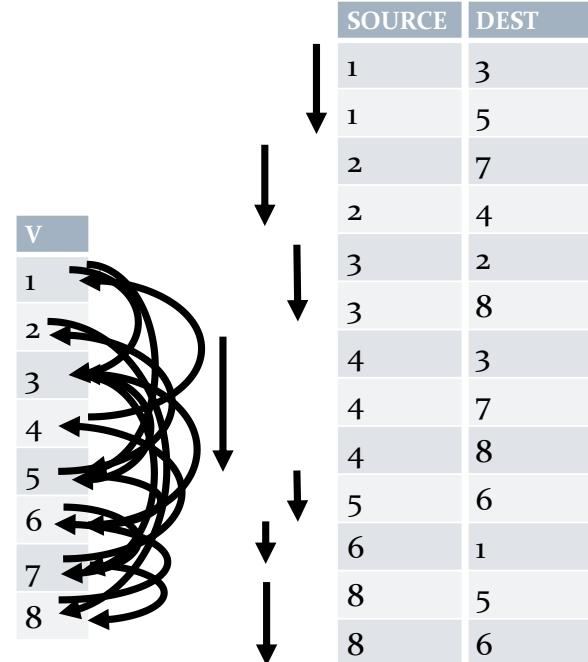


Random accesses over edges
Vertex-Centric Scatter Gather for BFS

-- "Think as an Edge" philosophy

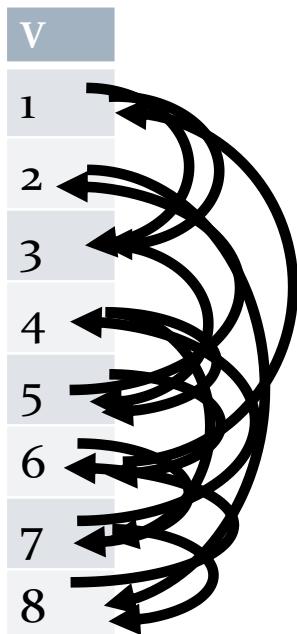
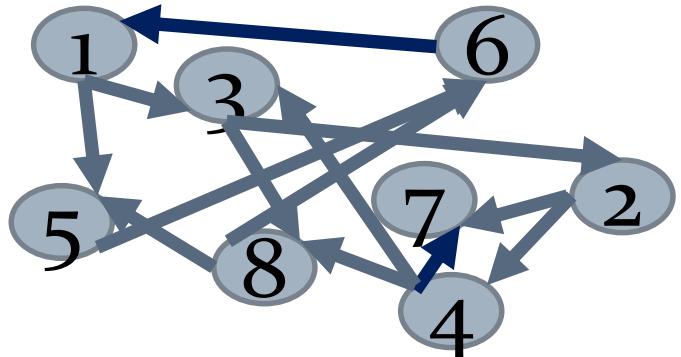


Transformation



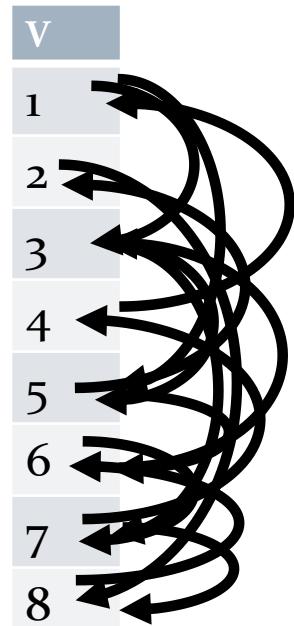
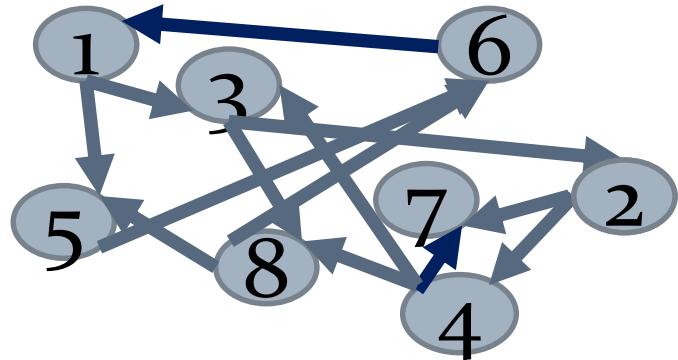
Sequential accesses over edges
Edge-Centric Scatter Gather for BFS

Illustration of vertex-centric BFS



SOURCE	DEST
1	3
1	5
2	7
2	4
3	2
3	8
4	3
4	7
4	8
5	6
6	1
8	5
8	6

Illustration of edge-centric BFS

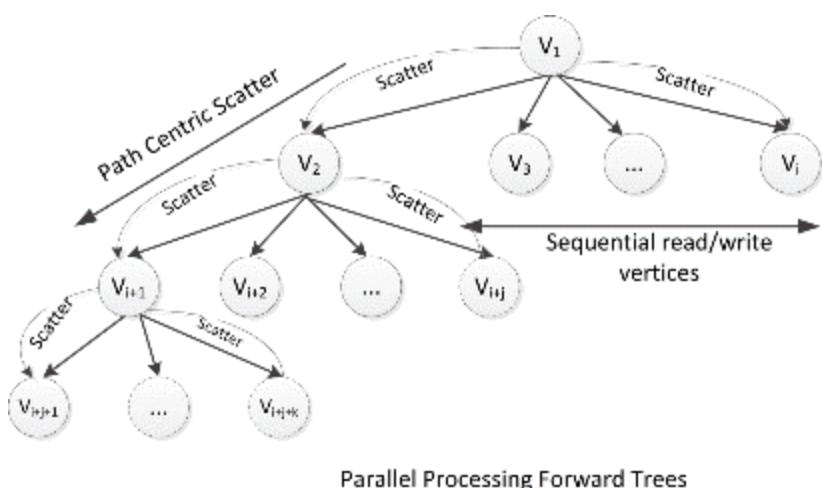


SOURCE	DEST
1	3
1	5
2	7
2	4
3	2
3	8
4	3
4	7
4	8
5	6
6	1
8	5
8	6

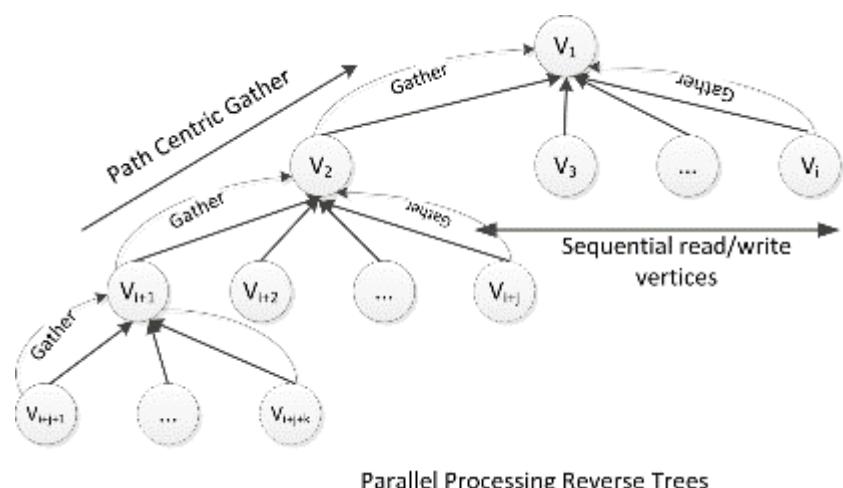
Programming model: path-centric

Computation model

- Scatter or gather: faster method
- Process graph following paths: improve locality



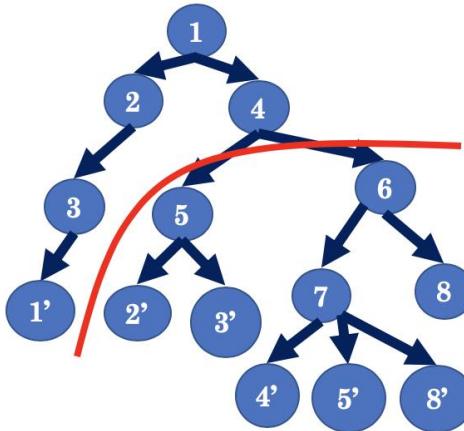
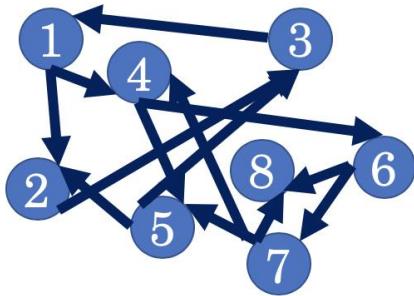
Path Centric Scatter



Path Centric Gather

Programming model: path-centric

- In order to achieve it, generate edge travel tree for each graph, then partition tree into multiple parts
 - edge travel tree: each edge is visited once. So, each branch of the tree is a path. The right graph is an edge traversal tree of the left graph.
- ◆ Each partition with forward-visited tree and backward-tree
- ◆ Each partition can parallelly executed



```

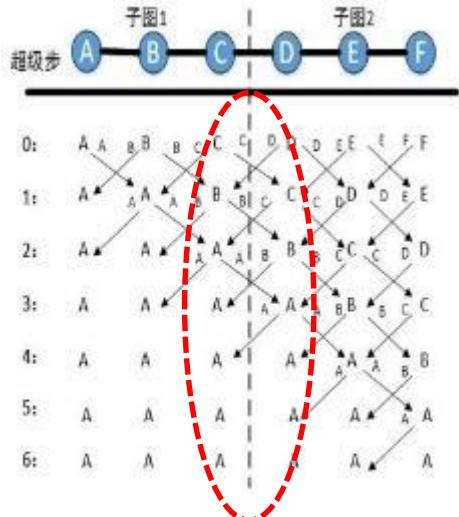
1: for each iteration do
2: parfor each  $p$  of Partitions
do
3:   Gather( $p$ );
4: end parfor
5: sync;
6: end for
  
```

Programming model: subgraph-centric

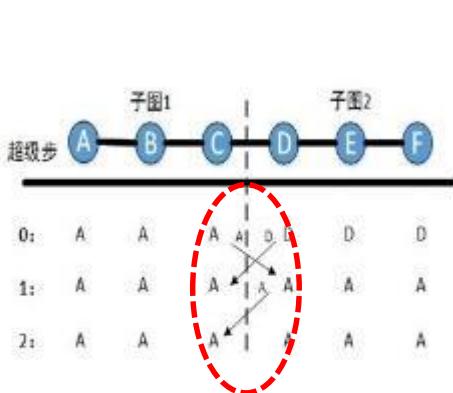
Coding graph algorithms as **subgraph-centric** programs to process vertices/edges via taking subgraph as the minimum parallel processing unit.

-- "Think as a Subgraph" philosophy

- ◆ First, it divides the graph data into different subgraphs
- ◆ Then, it updates all vertices of each subgraph until the subgraph converges
- ◆ Finally, it transmits the updated state information of the subgraph to other subgraphs



Many high cost global synchronization operations



Less global synchronization operations

Subgraph-centric Scatter Gather for CC

--It not only ensures better data locality, but also can transform many high cost global synchronization operations to low overhead local ones.

- The status information can only be propagated one-hop in each super-step for vertex/edge centric
- 7 super-step to converge for vertex/edge-centric
- 3 super-step to converge for subgraph-centric

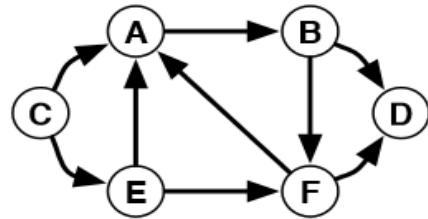
Programming model: subgraph-centric

Coding graph algorithms as **subgraph-centric** programs to process vertices/edges via taking subgraph as the minimum parallel processing unit.

-- "Think as a Subgraph" philosophy

- ◆ **First**, it divides the graph data into different subgraphs

examples of subgraphs



(a)

<u>Partition</u>	<u>Vertex</u>	<u>Edge List</u>	<u>Subgraph</u>
P1	(A)	[B]	G1
	(B)	[D F]	
P2	(C)	[A E]	G2
	(D)	[]	
P3	(E)	[A F]	G3
	(F)	[A D]	

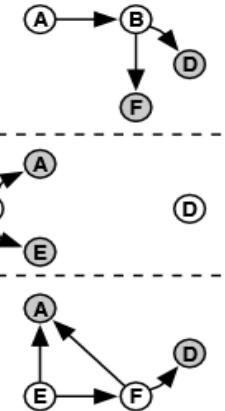


Figure (a) shows the original graph.

Figure (b) divides the set of vertices in the original graph into partitions.

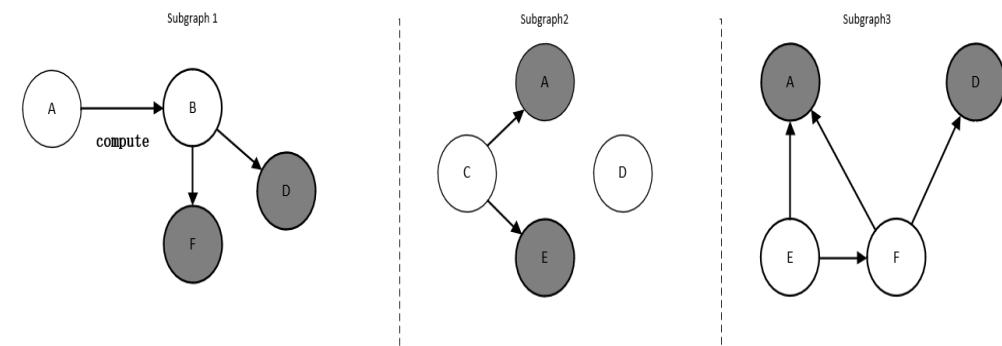
Figure (c) shows examples of subgraphs.

Programming model: subgraph-centric

Coding graph algorithms as **subgraph-centric** programs to process vertices/edges via taking subgraph as the minimum parallel processing unit.

-- "Think as a Subgraph" philosophy

- ◆ Then, it updates all vertices of each subgraph until the subgraph converges.



TLASG model

TLASG model

Compute(vertex v)
update all vertices of subgraph

```
block_update(subgraph subg)
while not done
for all vertices v that have updates
    compute(vertex v)
    apply updates from inbound edges of subgraph
```

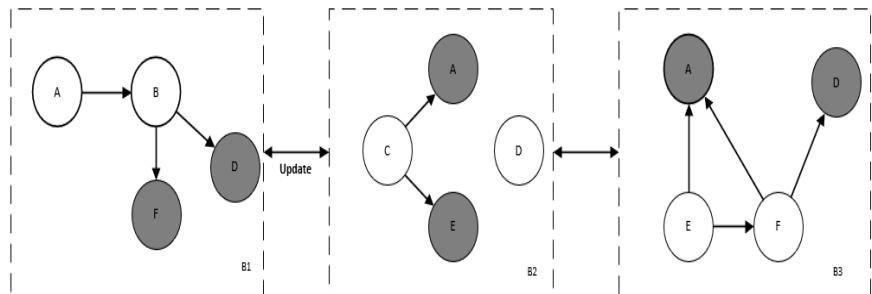
```
while not done
for all subg that have updates
    block_upddate(subg)
```

Programming model: subgraph-centric

Coding graph algorithms as **subgraph-centric** programs to process vertices/edges via taking subgraph as the minimum parallel processing unit.

-- "Think as a Subgraph" philosophy

- ◆ Finally, it transmits the updated state information of the subgraph to other subgraphs.



TLASG model

Compute(vertex v)
update all vertices of subgraph

```
block_update(subgraph subg)
  while not done
    for all vertices v that have updates
      compute(vertex v)
    apply updates from inbound edges of subgraph
```

```
while not done
  for all subg that have updates
    block_upddate(subg)
```

Graph processing frameworks

Programming Model

Vertex Centric: Pregel^[SIGMOD'10]
Edge Centric: X-Stream^[SOSP'13]
Path Centric: PathGraph^[SC'14]
Subgraph centric: Blogel^[VLDB'14]

Graph Partition Strategy

Vertex Cut: Pregel^[SIGMOD '10]
Edge cut: PowerGraph^[OSDI'12]
Hybrid Cut: PowerLyra^[EuroSys '15]

Computation Strategy

Pull vs push

Synchronous vs Asynchronous

Typical System

Pregel : SIGMOD 2010
PowerGraph: OSDI 2012
GraphChi: OSDI 2012

Graph partition

Graph is too large to hold in single node memory.

Goal of graph partition

- Balance workload between machines
- Minimize communication
- Maximize the computation efficiency within each machine

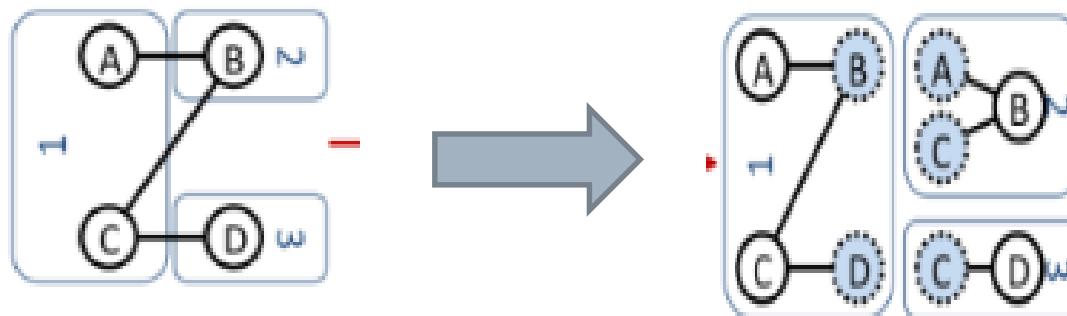
Graph partition method

- Edge-cut
- Vertex-cut
- Hybrid-cut

Graph partition: edge-cut

Pregel^[SIGMOD'08] and **GraphLab**^[VLDB'12]

- Focus on exploiting **locality**
- Partitioning: use **edge-cut** to evenly assign vertices along with all edges
- Computation: **aggregate** all resources (i.e., messages or replicas) of a vertex on local machine



Graph partition: edge-cut

Natural Graphs are Skewed

Graphs in real-world

- *Power-law* degree distribution
(aka Zipf's law or Pareto)

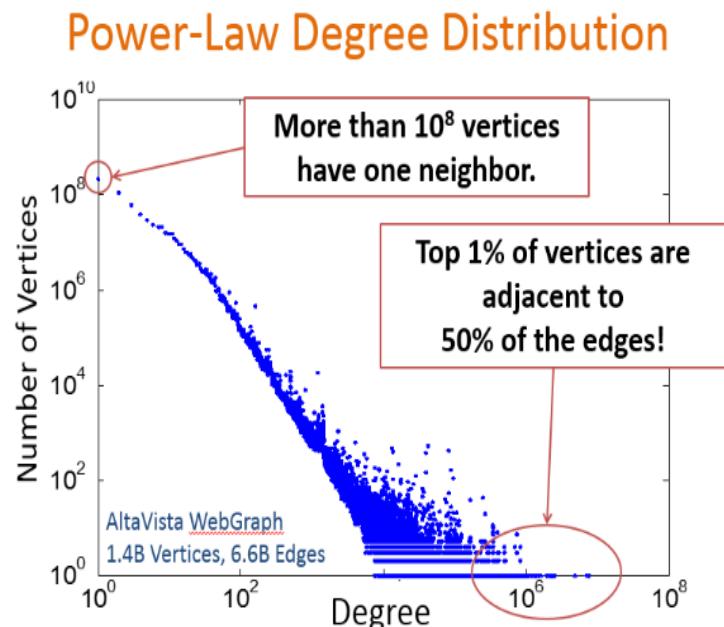
“**most** vertices have relatively **few** neighbors while a **few** have **many** neighbors” -- *PowerGraph*^[OSDI'12]

Case: SINA Weibo

- 167 million active users



#	User Profile	Follower Count
#1		77,971,093
#2		76,946,782
#100		14,723,818
...
397	...	397
...	...	69
...



Disadvantages of edge-cut

Load balance

- Work imbalance for high degree vertices, as computation, storage, and communication are linear with the degree of vertices

Difficult to partition

- Natural graphs are difficult to partition to minimize the communication, and maximize work balance

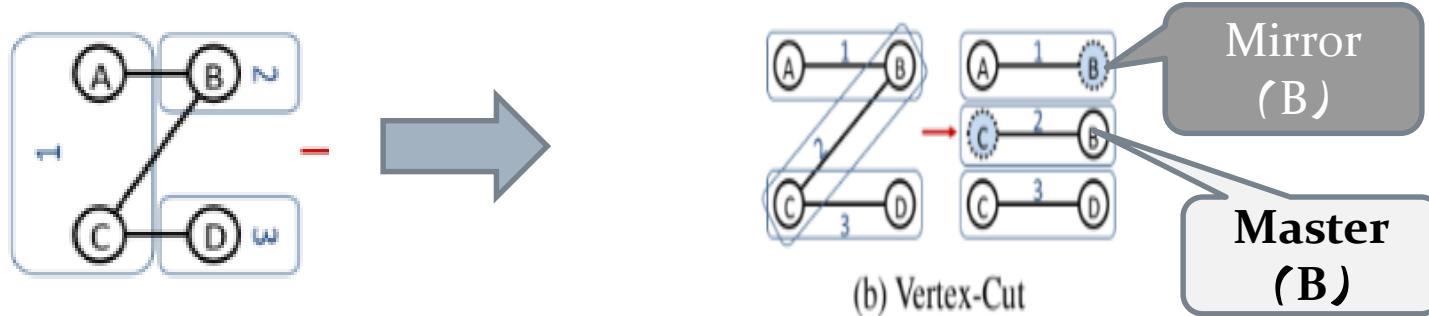
Limited parallelism

- No parallelism possible within individual vertices

Graph partition: vertex-cut

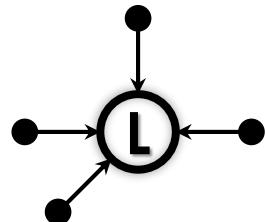
PowerGraph [OSDI'12] and **GraphX** [OSDI'14]

- Focus on exploiting **parallelism**
- Partitioning: use **vertex-cut** to evenly assign edges with replicated vertices
- Computation: **decompose** the workload of a vertex into multiple machines



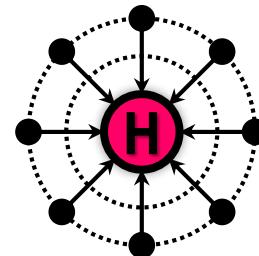
Is vertex-cut perfect for power-law graph?

Low-degree vertex



397

High-degree vertex



77,971,093

Locality

make resource **locally** accessible to hidden network latency



100-million users each with 100 followers

Parallelism

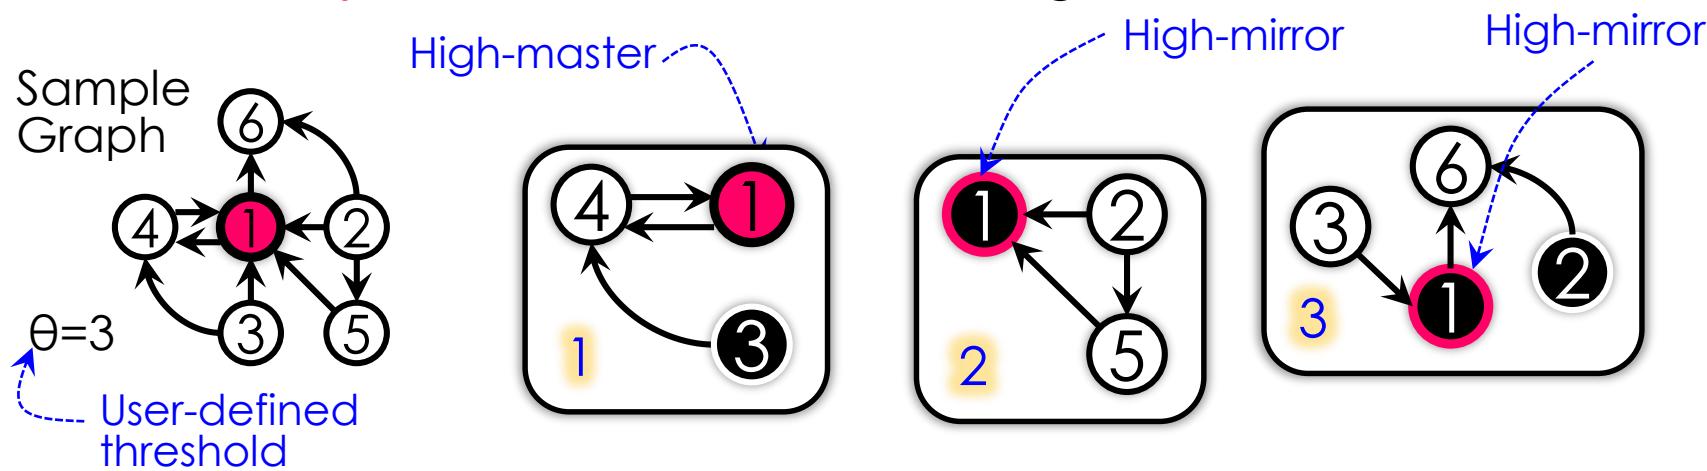
evenly **parallelize** the workloads to avoid load imbalance

100 users each with 100-million followers

Graph partition: hybrid-cut

Hybrid-cut: differentiate the graph partitioning for **low-degree** and **high-degree** vertices

- **Low-cut** (inspired by *edge-cut*): reduce mirrors and exploit locality for low-degree vertices
- **High-cut** (inspired by *vertex-cut*): provide balance and restrict the impact of high-degree vertices
- **Efficiently** combine low-cut and high-cut



Graph processing frameworks

Programming Model

Vertex Centric: Pregel^[SIGMOD'10]
Edge Centric: X-Stream^[SOSP'13]
Path Centric: PathGraph^[SC'14]
Subgraph centric: Blogel^[VLDB'14]

Graph Partition Strategy

Vertex Cut: Pregel^[SIGMOD '10]
Edge cut: PowerGraph^[OSDI'12]
Hybrid Cut: PowerLyra^[EuroSys '15]

Computation Strategy

Pull vs push

Synchronous vs Asynchronous

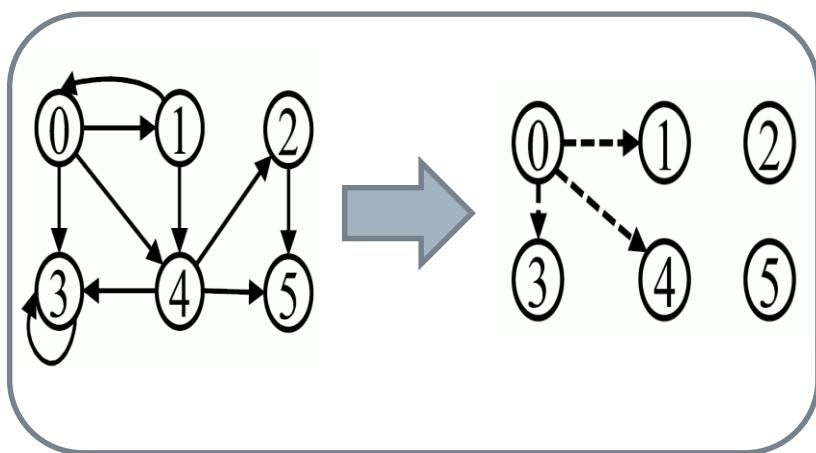
Typical System

Pregel : SIGMOD 2010
PowerGraph: OSDI 2012
GraphChi: OSDI 2012

Message passing: push model

Push updates to the neighbors

- Selectively schedule to reduce computation
- Lead to the write-write conflict on a destination vertex from multiple sources



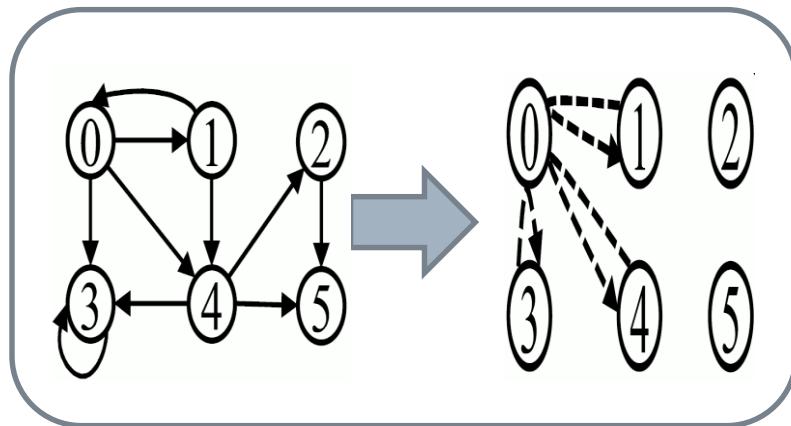
```
parallel_for (int vSrc = 0; vSrc < numVertices; ++vSrc) {  
    if (!frontier.contains(vSrc)) continue;  
    for (int d = 0; d < vertex[vSrc].outdegree; ++d) {  
        const int vDst = vertex[vSrc].outneighbor[d];  
        if (converged.contains(vDst)) continue;  
        atomicCAS(vertex[vDst].value,  
                  compute(vertex[vSrc].value, vertex[vDst].value)); } }
```

Push-based message implementation

Message passing: pull model

Get update from incoming neighbors

- Do not have the write-write conflict
- Scan all incoming neighbors can incur redundant computation since some nodes may not active

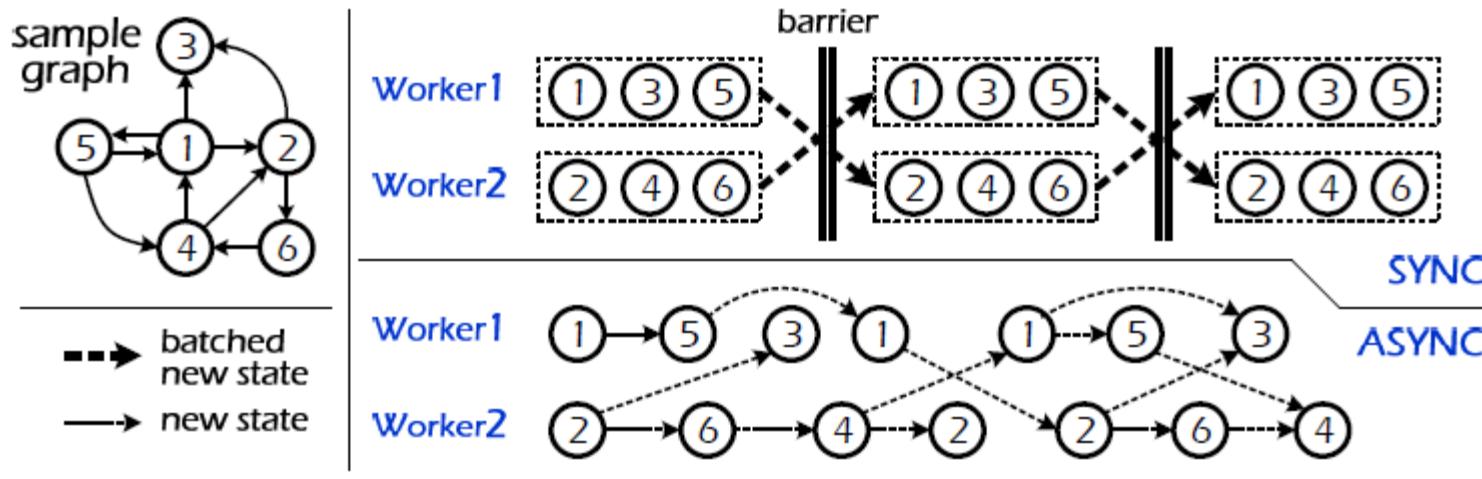


```
parallel_for (int vDst = 0; vDst < numVertices; ++vDst) {  
    if (converged.contains(vDst)) continue;  
    for (int s = 0; s < vertex[vDst].indegree; ++s) {  
        const int vSrc = vertex[vDst].inneighbor[s];  
        if (!frontier.contains(vSrc)) continue;  
        vertex[vDst].value =  
            compute(vertex[vSrc].value, vertex[vDst].value); } }
```

Pull-based message implementation

Schedule model

Synchronous vs Asynchronous

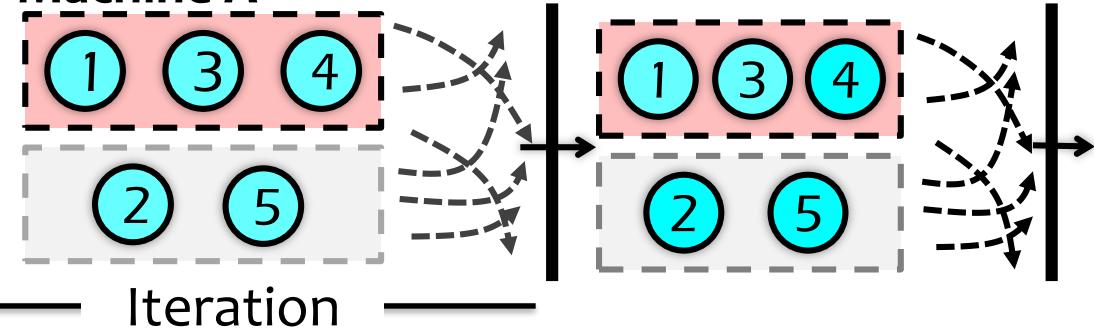


	Sync	Async
Properties		
Communication	Regular	Irregular
Convergence	Slow	Fast
Favorites		
Algorithm	I/O-intensive	CPU-intensive
Execution Stage	High Workload	Low Workload
Scalability	Graph Size	Cluster Size

Schedule model: sync model

Scheduling:

Machine A



Current iteration active vertex

next iteration active vertex

Pseudocode

while ($\text{iteration} \leq \text{max}$) **do**

if $V_a == \emptyset$ **then break**
 $V'_a \leftarrow \emptyset$

foreach $v \in V_a$ **do**
 $A \leftarrow \text{compute}(v)$
 $V'_a \leftarrow V'_a \cup A$

barrier to update
 $V_a \leftarrow V'_a$

iteration ++

Ref. PoweSwitch PPoPP 2015

Schedule model: sync model

Advantages of synchronous model

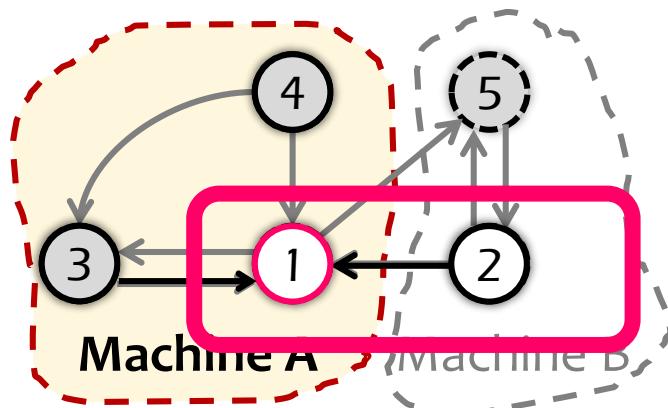
1. Batching messages which improves network utilization
2. Better for algorithm with much communication needed
3. With light computation for each node

Disadvantages of synchronous model

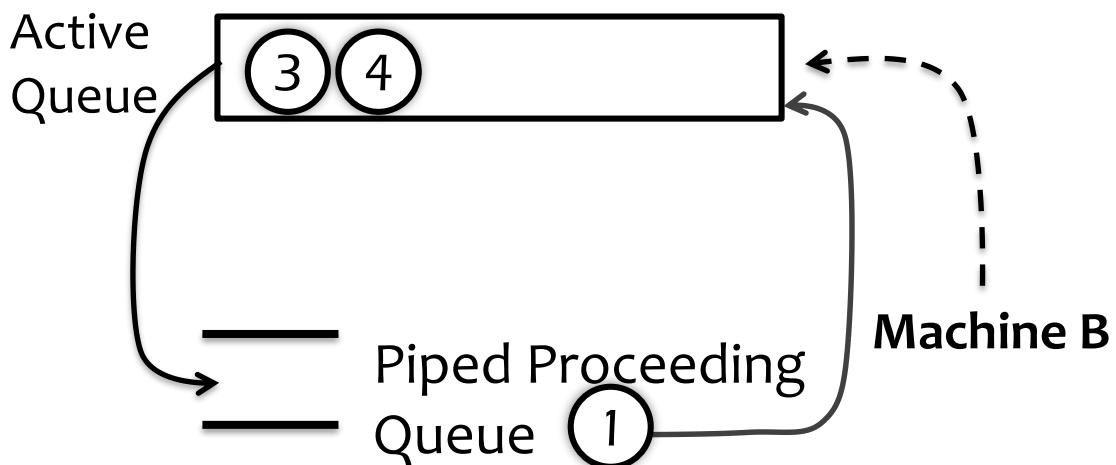
1. Uneven convergence problem
2. Straggler problem
3. Not suitable for some algorithms, eg. Graph Coloring and Clustering based on Gibbs Sampling

Schedule model: async model

Scheduling:



Internal State (e.g. Machine A):



Pseudocode

```

while ( $V_a \neq \emptyset$ ) do
     $v = \text{dequeue}(V_a)$ 
     $A \leftarrow \text{compute}(v)$ 
     $V'_a \leftarrow V'_a \cup A$ 
    signal across machines
  
```

Schedule model: async model

Advantages of asynchronous

1. Speedup convergence
2. Better for CPU-bound algorithm

Disadvantages of asynchronous

1. Lock cost caused by vertex contention
2. Asynchronous model sends messages frequently, and CPU frequently encapsulates TCP/IP packets, which wastes the utilization of CPU

Graph processing frameworks

Programming Model

Vertex Centric: Pregel^[SIGMOD'10]
Edge Centric: X-Stream^[SOSP'13]
Path Centric: PathGraph^[SC'14]
Subgraph centric: Blogel^[VLDB'14]

Graph Partition Strategy

Vertex Cut: Pregel^[SIGMOD '10]
Edge cut: PowerGraph^[OSDI'12]
Hybrid Cut: PowerLyra^[EuroSys '15]

Computation Strategy

Pull vs push
Synchronous vs Asynchronous

Typical System

Pregel : SIGMOD 2010
PowerGraph: OSDI 2012
GraphChi: OSDI 2012

Graph processing system

□ In-memory single-node graph system

- Ligra, GraphMat, Polymer

□ Out-of core single-node graph system

- GraphChi, TurboGraph, X-Stream, PathGraph, GridGraph, FlashGraph

□ Distributed graph system

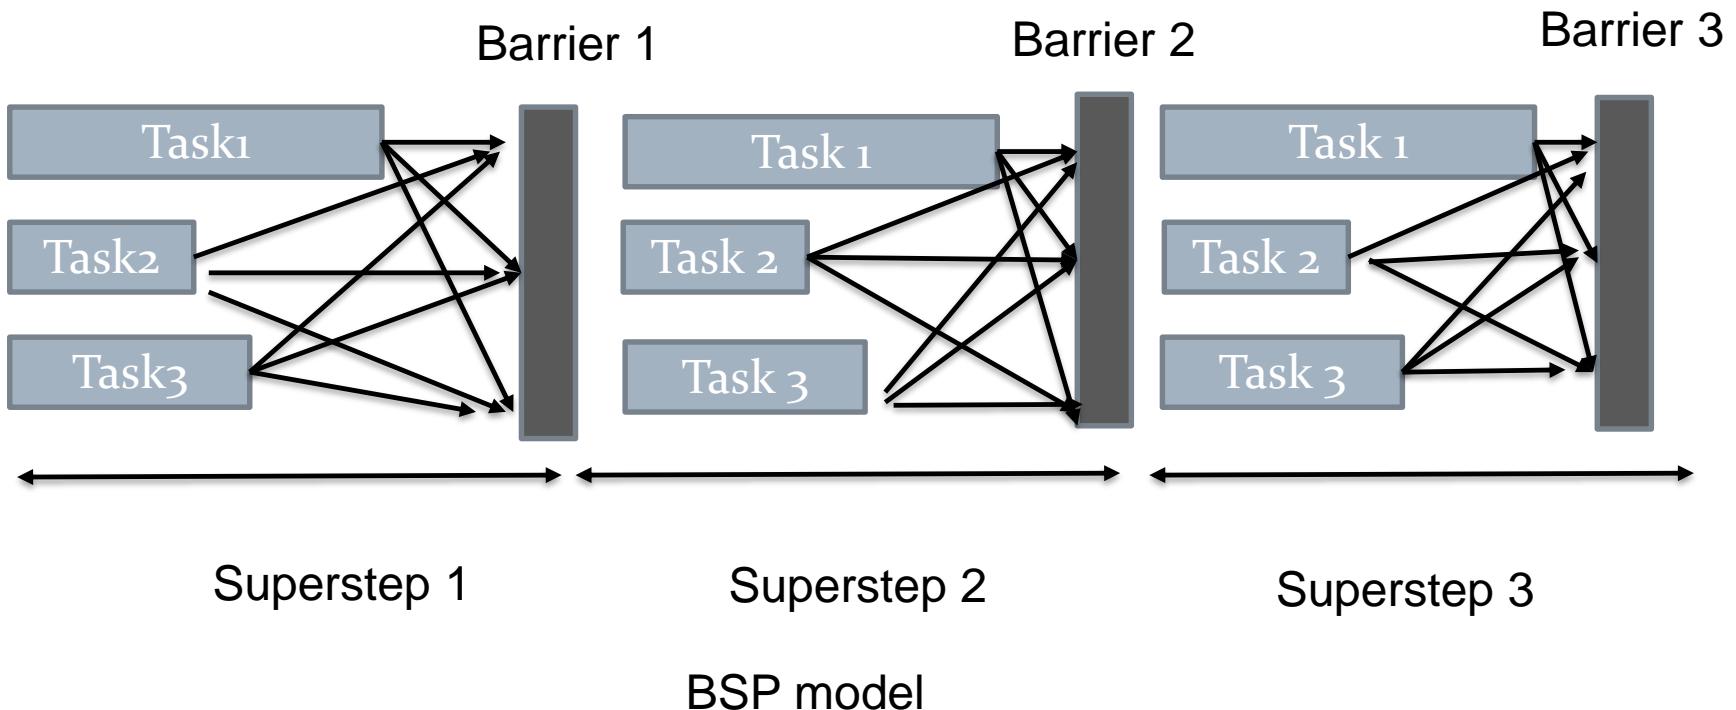
- Pregel, PowerGraph, GraphLab, GraphX, PowerSwitch, PowerLyra, Gemini

□ Graph system on accelerator

- GPU based, e.g. CuSha, WS, Gunrock, IrGL, Groute
- FPGA based
- ASIC based
- NVM based

Pregel

- Large scale graph-parallel processing platform developed by Google
- Utilize **BSP** (bulk synchronous parallel) model



Pregel

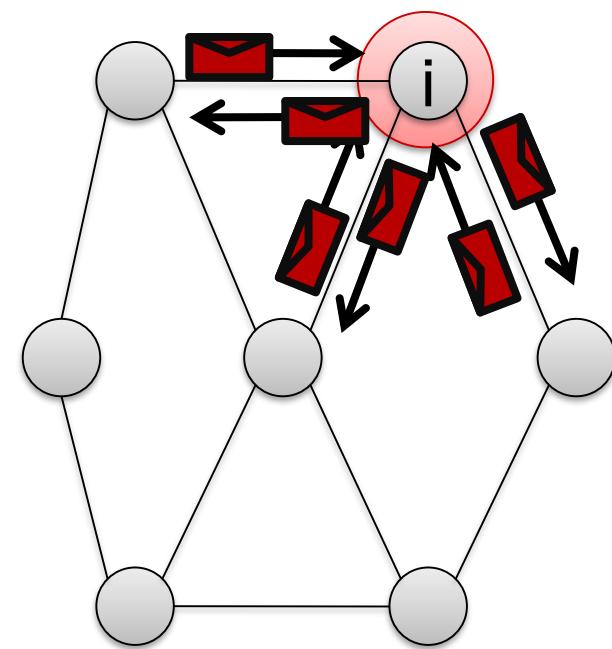
- Large scale graph-parallel processing platform developed by Google
- Utilize BSP (bulk synchronous parallel) model
- Vertex-Programs interact by sending messages

```
Pregel_PageRank(i, messages) :
```

```
// Receive all the messages
total = 0
foreach( msg in messages) :
    total = total + msg
```

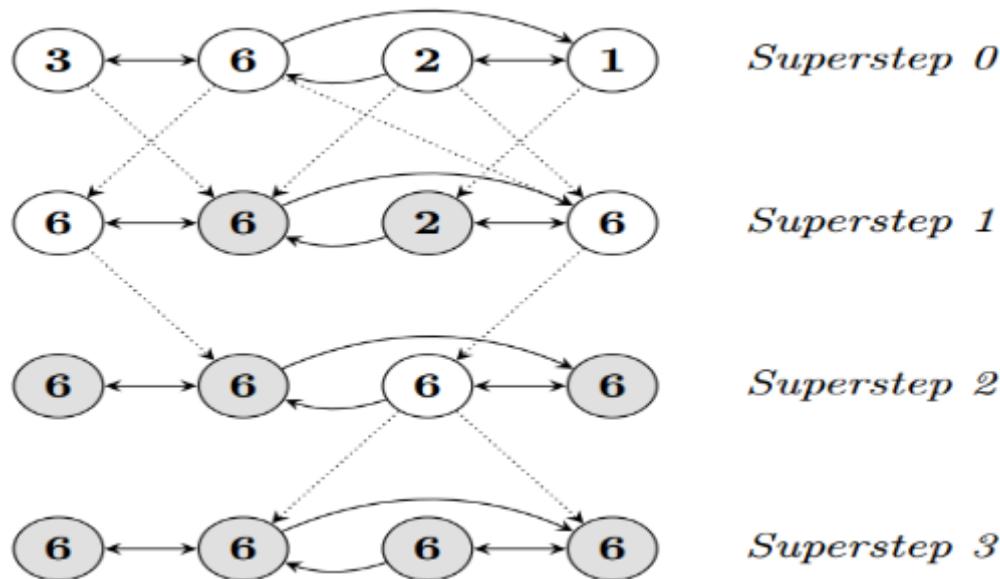
```
// Update the rank of this vertex
R[i] = 0.15 + total
```

```
// Send new messages to neighbors
foreach(j in out_neighbors[i]) :
    Send msg(R[i] * wij) to vertex j
```



Pregel

- Large scale graph-parallel processing platform developed by Google
- Utilize BSP (bulk synchronous parallel) model
- Vertex-Programs interact by sending messages.
- Iteration until convergence



Maximum value example for Pregel

PowerGraph

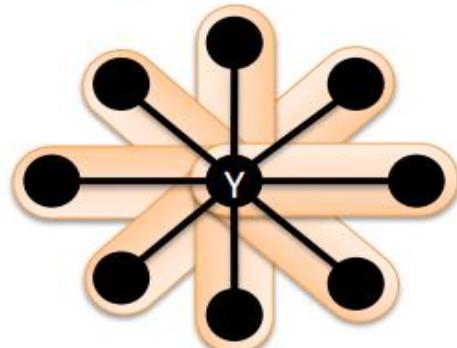
Introduce GAS programming model

Gather (Reduce)

Accumulate information about neighborhood

User Defined:

- ▶ **Gather**(→ Σ)
- ▶ $\Sigma_1 + \Sigma_2 \rightarrow \Sigma_3$



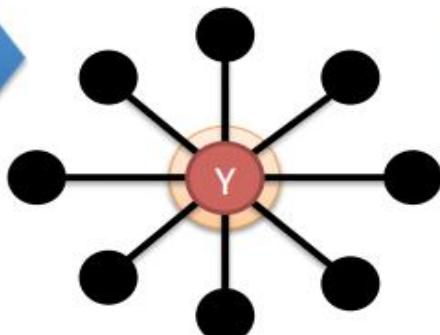
$$\text{Parallel Sum } \gg + \gg + \dots + \gg \rightarrow \Sigma$$

Apply

Apply the accumulated value to center vertex

User Defined:

- ▶ **Apply**(, Σ) → Σ'

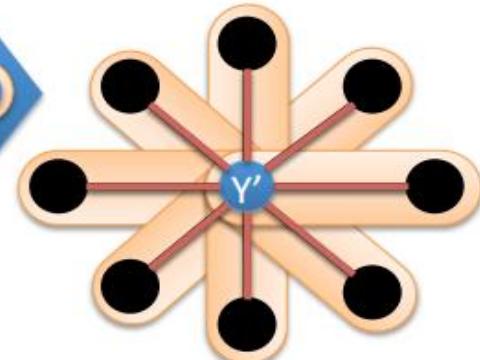


Scatter

Update adjacent edges and vertices.

User Defined:

- ▶ **Scatter**(→ —)

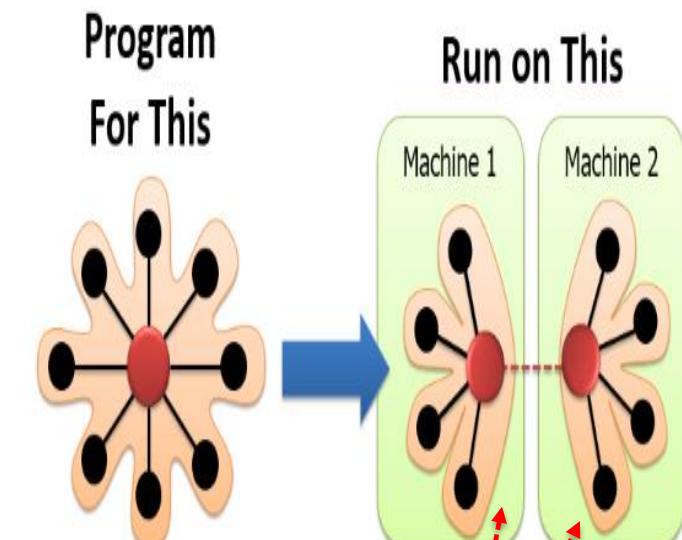


Update Edge Data & Activate Neighbors

33

PowerGraph

Edge-cut for power-law graph



Split High-Degree vertices

Cut high-degree vertex can balance workload

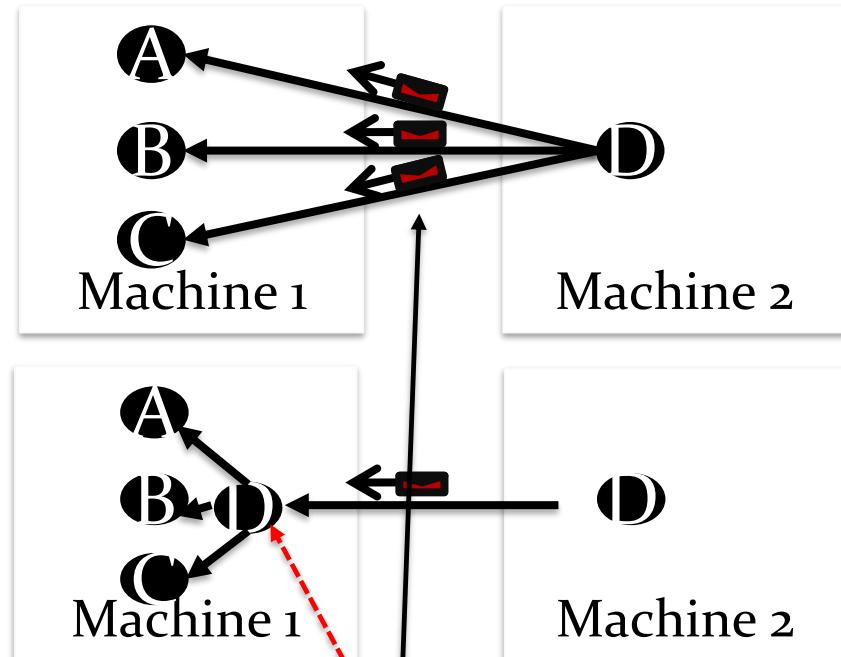


Image node can reduce communication between node

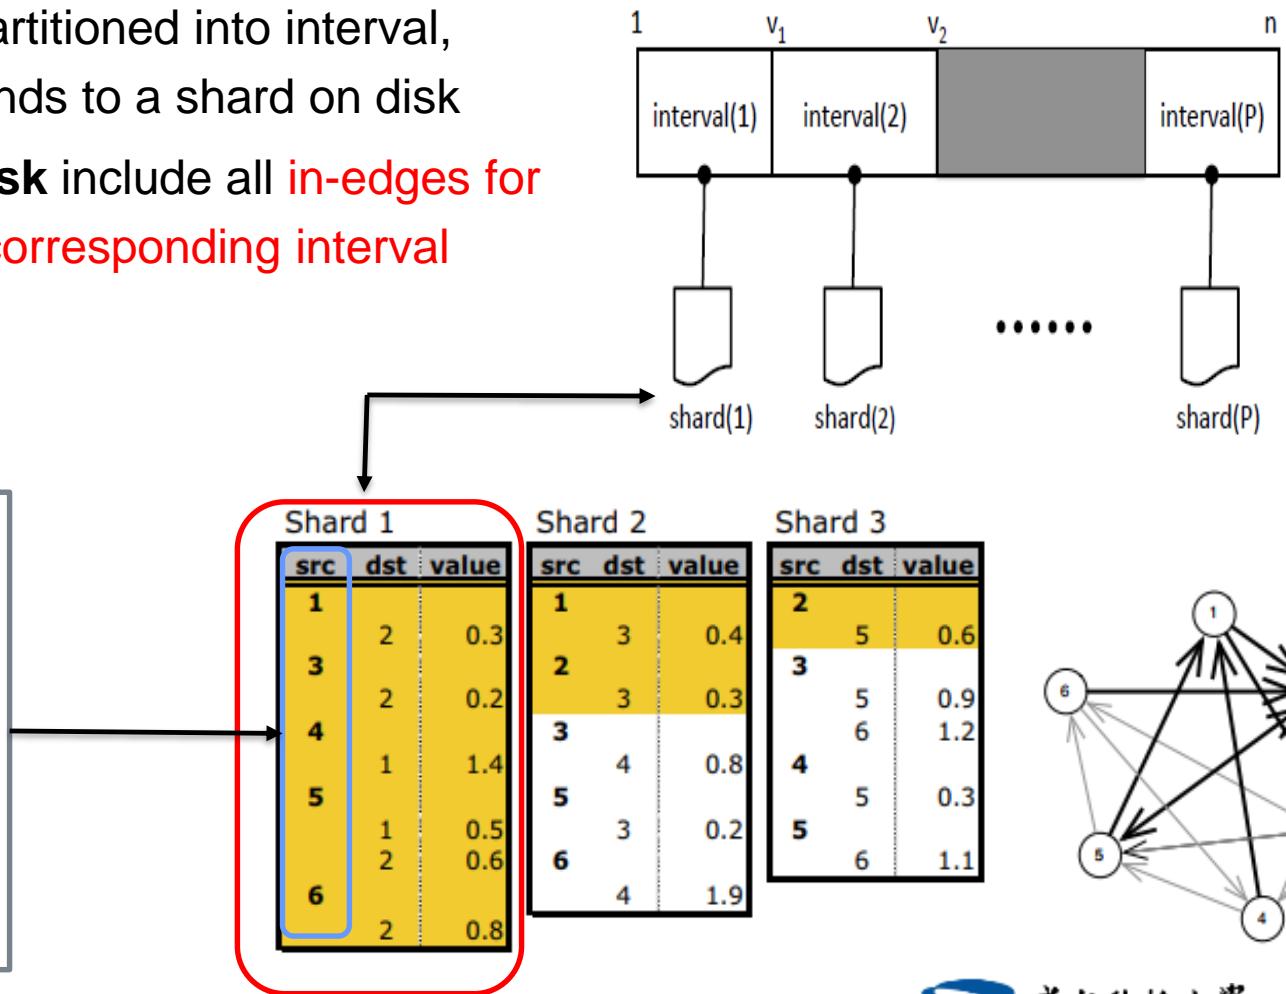
GraphChi

- Challenges for distributed graph computation
 - Partitioning a graph is difficult (especially for power law graphs)
 - Distributed graph processing incurs too much communication cost
- Could we compute big graphs on a single machine?
- **GraphChi** : Single-node out-of-core graph processing system
 - **Goal:** maximize sequential disk access while loading graph into memory (500x speedup for disk sequential vs random)
 - **Method:** Execute on individual subgraphs once a time with parallel sliding window (PSW) , load subgraphs efficiently from disk

GraphChi

PSW: parallel sliding window

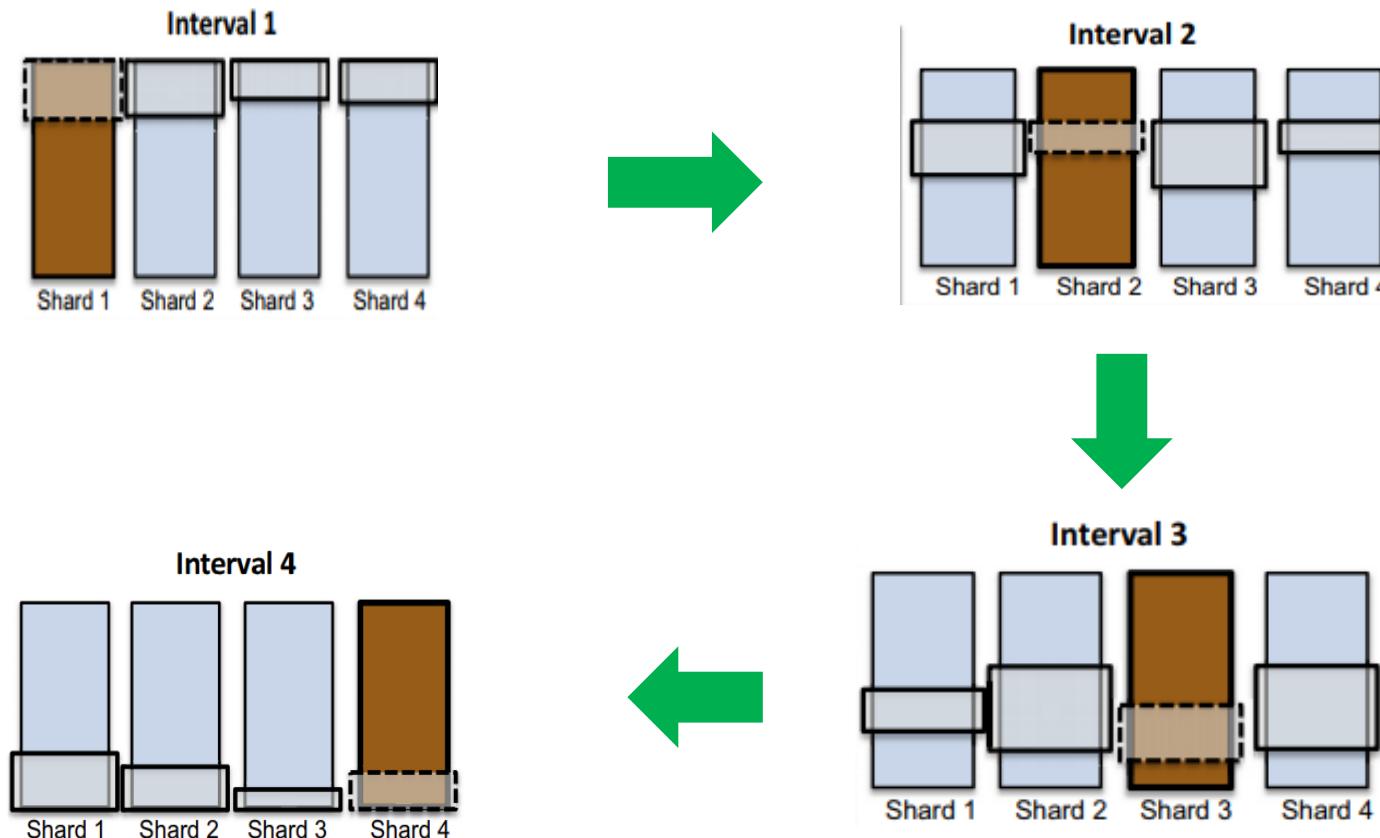
- Vertices are partitioned into interval, each corresponds to a shard on disk
- A shard on disk include all in-edges for all vertices in corresponding interval**



- Edges in a shard are sorted by source vertices
- Bring benefit for date access from disk
- Incur much preprocessing cost

GraphChi

- Slide window to process one intervals with shard one by one



GraphChi

- Processing workflow for one interval
(eg. Processing Interval 1)

Step 1: Load all in-edges from shard 1 into memory for vertices of interval 1

(1 sequential disk read)

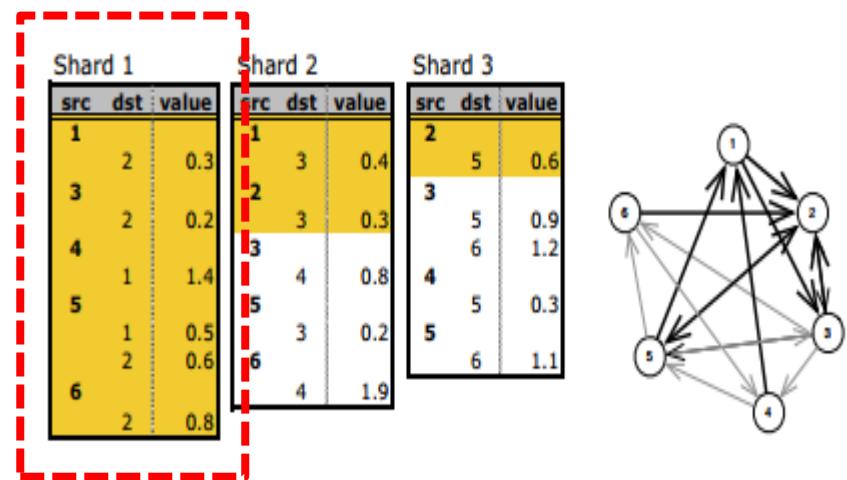


Fig. Execution interval 1 (with vertices 1, 2)

GraphChi

- Processing workflow for one interval
(eg. Processing Interval 1)

Step 1: Load all in-edges from shard 1 into memory for vertices of interval 1

Step 2: Load all out-edges in memory
form other shards for vertices of interval 1
(P-1 sequential reads)

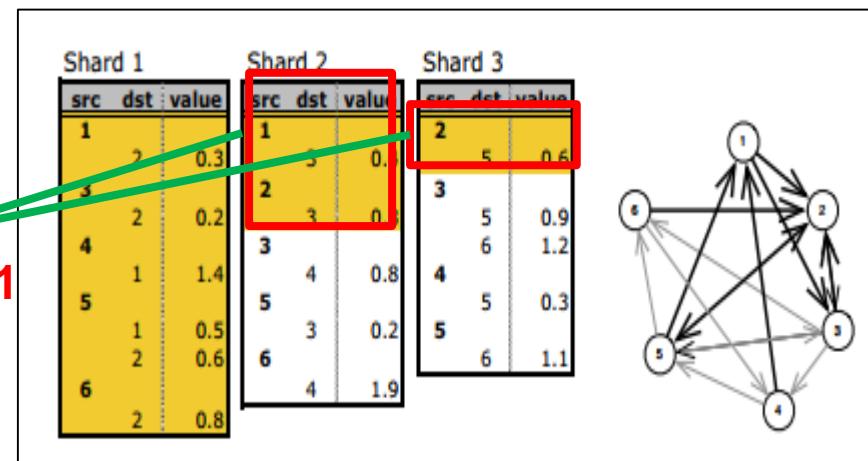


Fig. Execution interval (vertices 1, 2)

GraphChi

- Processing workflow for one interval
(eg. **Processing Interval 1**)

Step 1: Load all in-edges from shard 1 into memory for vertices of interval 1

Step 2: Load all out-edges in memory from other shard for vertices of interval 1

Step 3: Update-function is executed on interval's vertices (vertex 1, 2).

(vertices-centric processing)

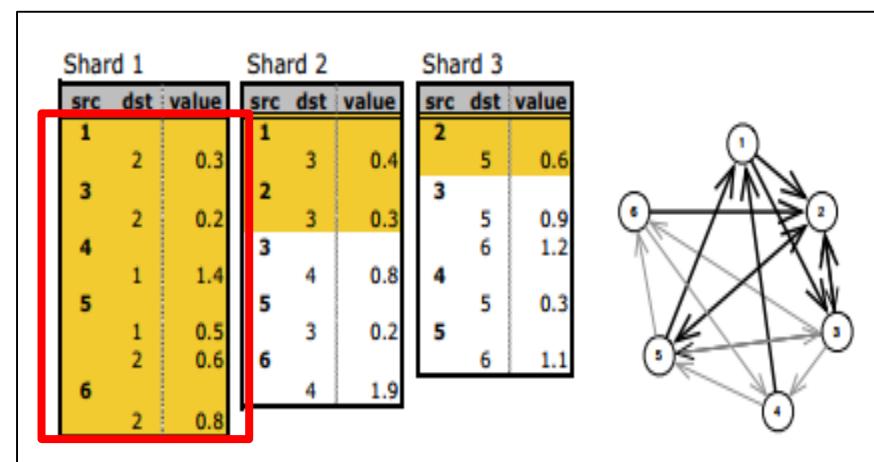


Fig. Execution interval 1 (with vertices 1, 2)

GraphChi

- Processing workflow for one interval
(eg. **Processing Interval 1**)

Step 1: Load all in-edges from shard 1 into memory for vertices of interval 1

Step 2: Load all out-edges in memory from other shard for vertices of interval 1

Step 3: Update-function is executed on interval's vertices (vertex 1, 2)

**Step 4: Commit to disk, updates are written back to disk
(p sequential disk write)**

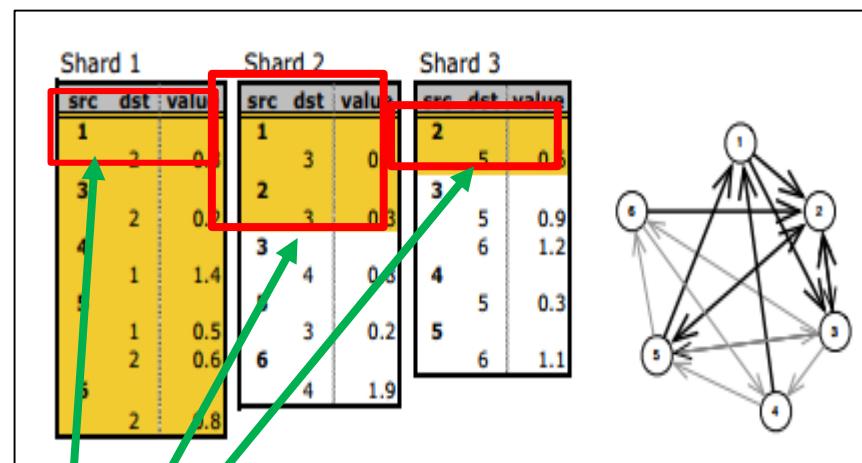


Fig. Execution interval 1 (with vertices 1, 2)

GraphChi

- Processing workflow for one interval
(eg. **Processing Interval 2**)

Step 1: Load all in-edges from shard 2 into memory for vertices of interval 2

Step 2: Load all out-edges in memory from other shard for vertices of interval 2

Step 3: Update-function is executed on interval's vertices (vertex 3, 4)

Step 4: Commit to disk, update are written back to disk
(p sequential disk write)

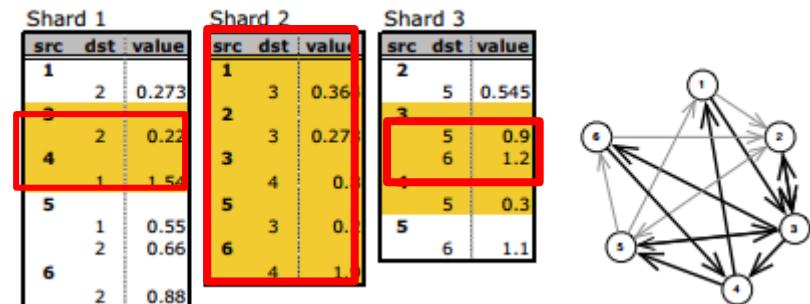


Fig. Execution interval 2 (vertices 3, 4)

PROGRAMMING FOR STREAM PROCESSING

Outline

- Stream data and stream applications
- Real-time stream processing
- Ecosystems and programming models

What is stream?



What is stream?

- A **stream** is a **continuous and unbounded sequence of data elements** made available over time
- Stream data in big data era
 - Volume, continuous and unbounded
 - Variety, changes quickly
 - Velocity, high speed incoming
 - Value, needs to be processed in near real-time



Stream applications



“ 11 - 11 ”



(Didi) Real-time Traffic Monitoring

热搜榜		更多 >
热	关注重庆公交坠江事件	热
1	金庸去世	爆 480万
2	这两天怎么了	新 270万
3	超会挑联盟	荐 264万
4	王光英去世	新 249万
5	唐嫣罗晋婚礼主持	热 176万
6	香港四大才子只剩蔡澜和...	新 139万
7	陈小春悼念金庸	新 122万
8	公交坠江发现9名遇难者	热 75万
9	悲伤的十月	新 72万
10	哈文朋友圈 好友辟谣	热 71万

Hot Search Topics Ranking



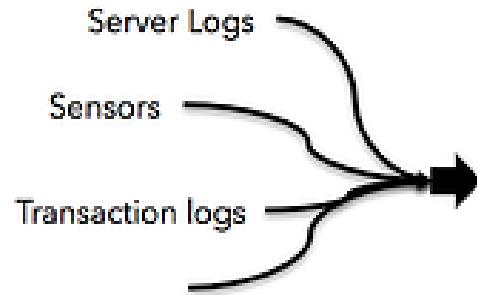
Market Transaction

Outline

- Stream data and stream applications
- Real-time stream processing
- Ecosystems and programming models

Common stream architecture

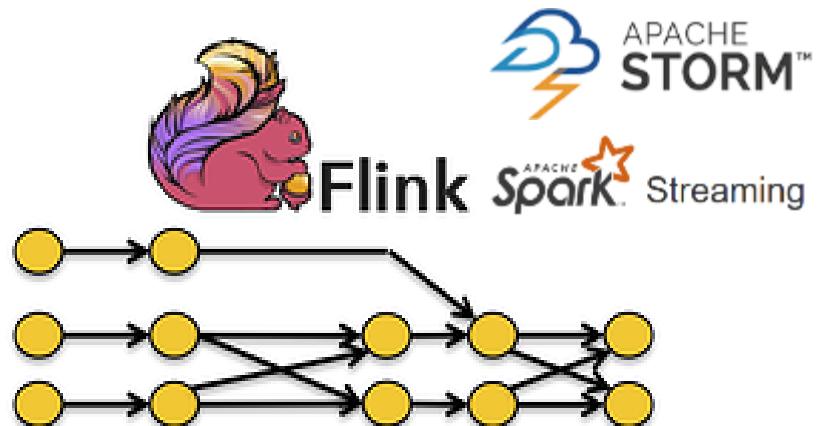
- Collect data from a variety of data sources
- Make data streams available for consumption
- Real-time stream processing



Gathering

Broker

Analysis



Data model - Tuples

- A tuple is the basic data element that can be processed independently
- Each tuple can be represented as a list of values. ($t = <value_1, value_2, value_3, \dots>$)

The image shows a user interface for posting a status update. The text input field contains "Hello world!". Below the input field are several icons: smiley face (表情), image (图片), video (视频), hash tag (话题), lightning bolt (头条文章), and three dots (...). To the right of the input field are buttons for "公开" (Public) and "发布" (Post). Above the input field, there is a placeholder text "有什么新鲜事想告诉大家？" and a character count indicator "已输入6字".

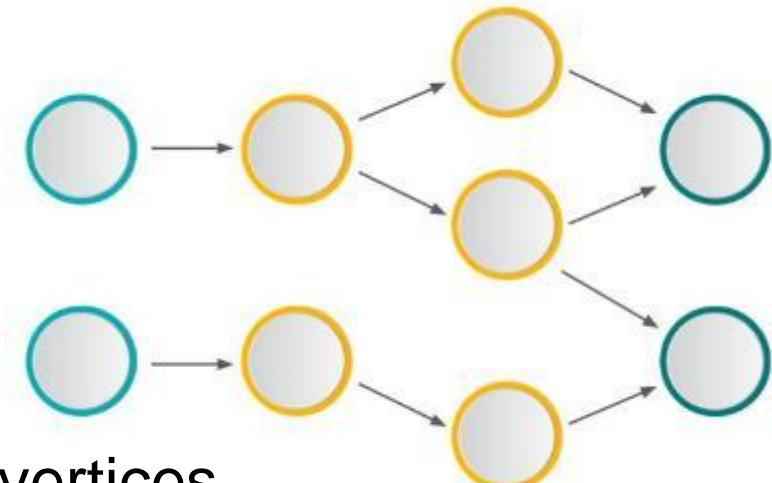
```
{  
    "user_id": 17055506,  
    "timestamp": 1421777578,  
    "status": "Hello world!"  
}
```

Application model - DAG

- Stream application can be modeled as a **Directed Acyclic Graph (DAG)**

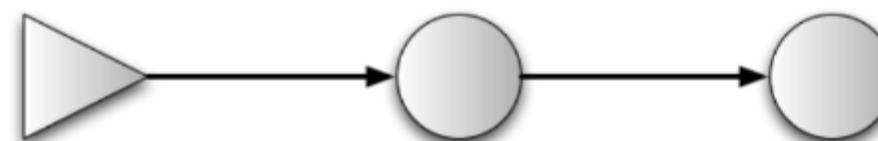
- Vertex: a simple **operator**

executes common user defined function, e.g., map, filter, reduce, join,...



- Edge: routing **tuples** between vertices

- A simple example: Twitter Word Count



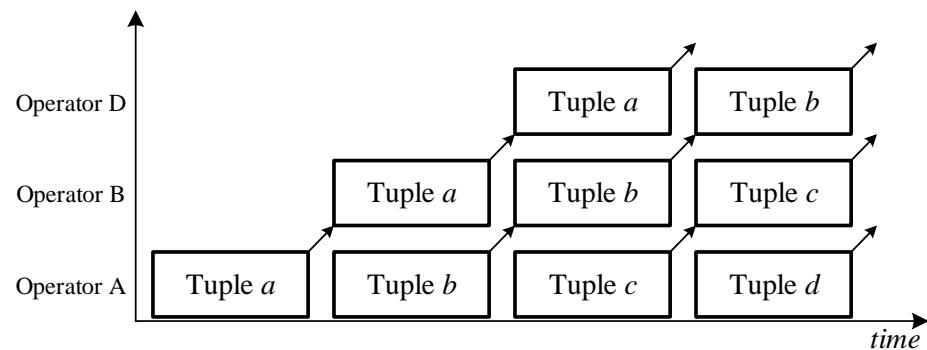
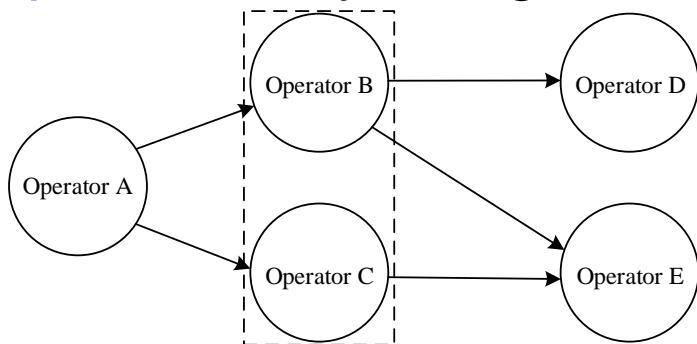
TweetSpout

ParseTweetBolt

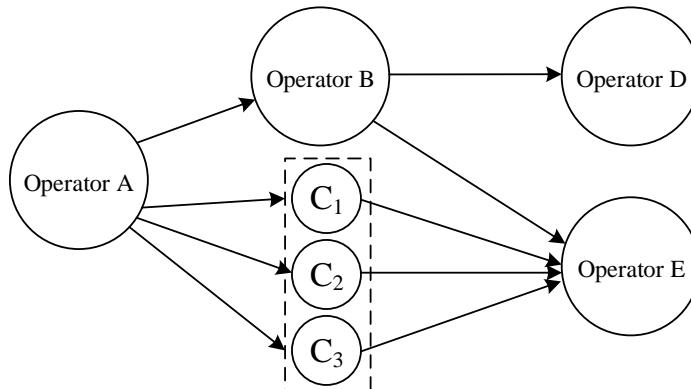
WordCountBolt

Parallisms in stream processing

- Stream processing exploits **pipeline parallelism** and **task parallelism** by using **DAG model**



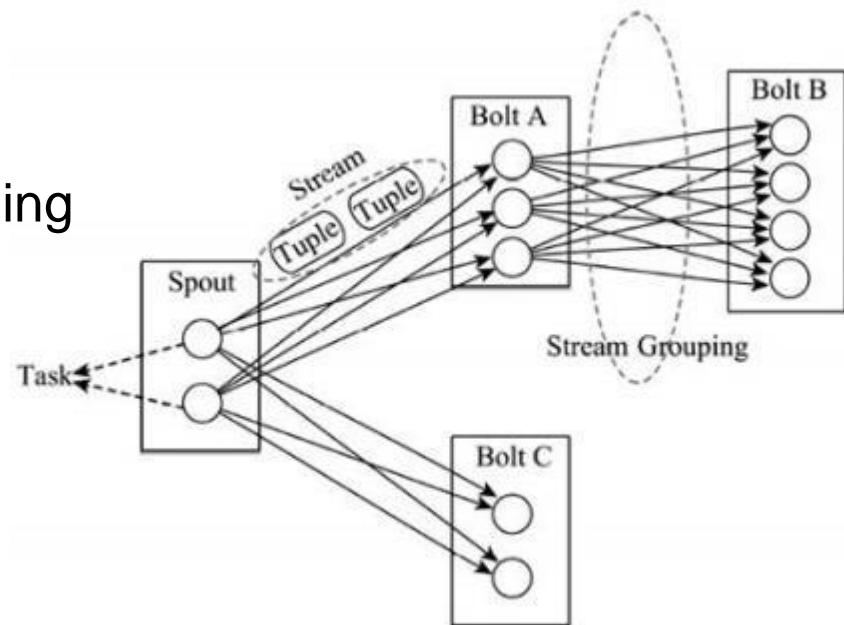
- Stream processing system can generate multiple **instances** for each operator to achieve high **data parallelism**



Tuples grouping and dispatching

- Tuples are partitioned into different **groups** and dispatched to different **instances** of operators

- Grouping strategies
 - Key Grouping: hash based partitioning
 - Shuffle Grouping: round-robin
 - All Grouping: broadcast
 - Global Grouping: all to one
 - Others



Programming: step by step

1. Read input streams from data sources and generate a sequence of **tuples**

- Kafka, HDFS, Flume, ...

2. Build the **DAG** for a stream application

3. Specify the **degree of parallelisms** for each operator and the data **grouping strategies**

- Depending on different platform APIs

4. Submit the program to the cluster

- Deployment
- Fault Tolerant
- Dashboard
- ...

Managed by platforms

Transparent to programmers

Outline

- Stream data and stream applications
- Real-time stream processing
- Ecosystems and programming models

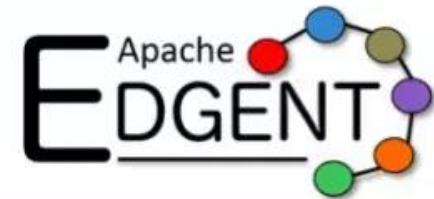
Ecosystems

Open Source Stream Processing Frameworks



Samza

Apache Flume™



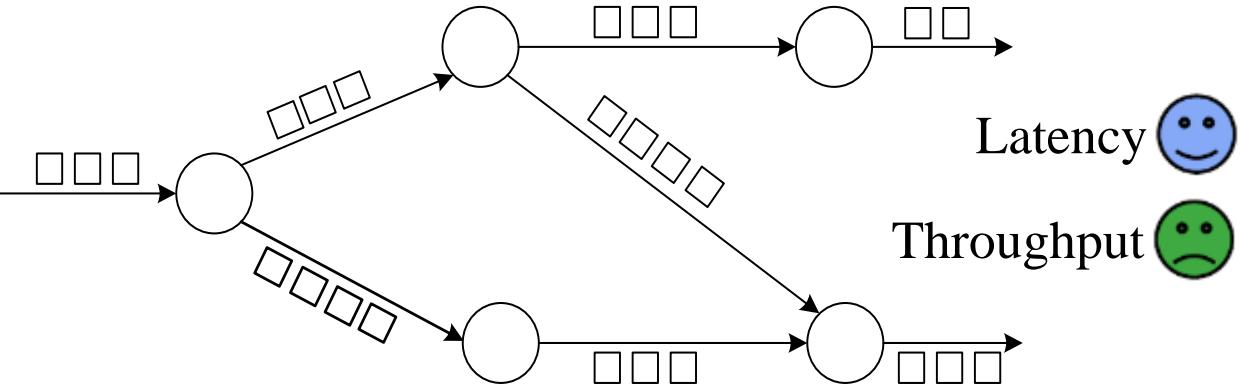
Apache Gearpump



Two different programming models

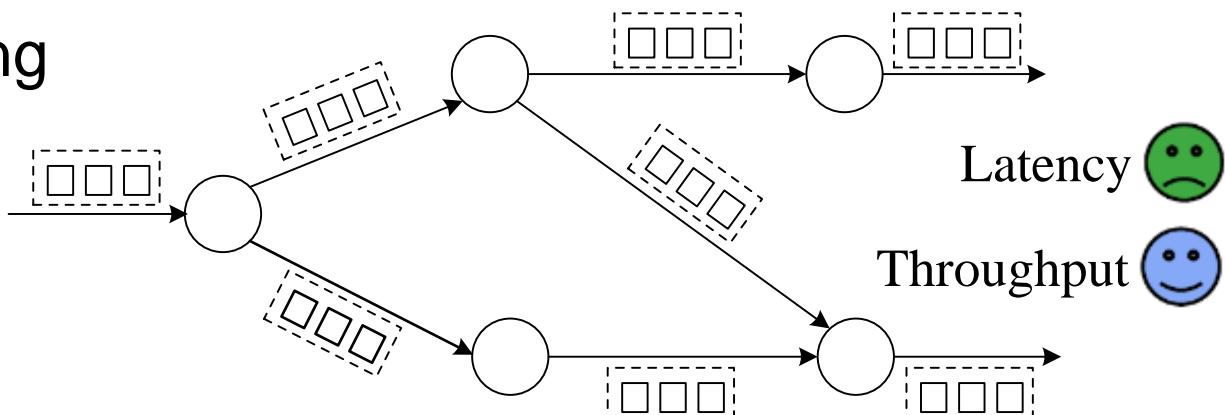
- Continuous dataflow: process tuples one by one

- Apache Storm
 - Apache Flink



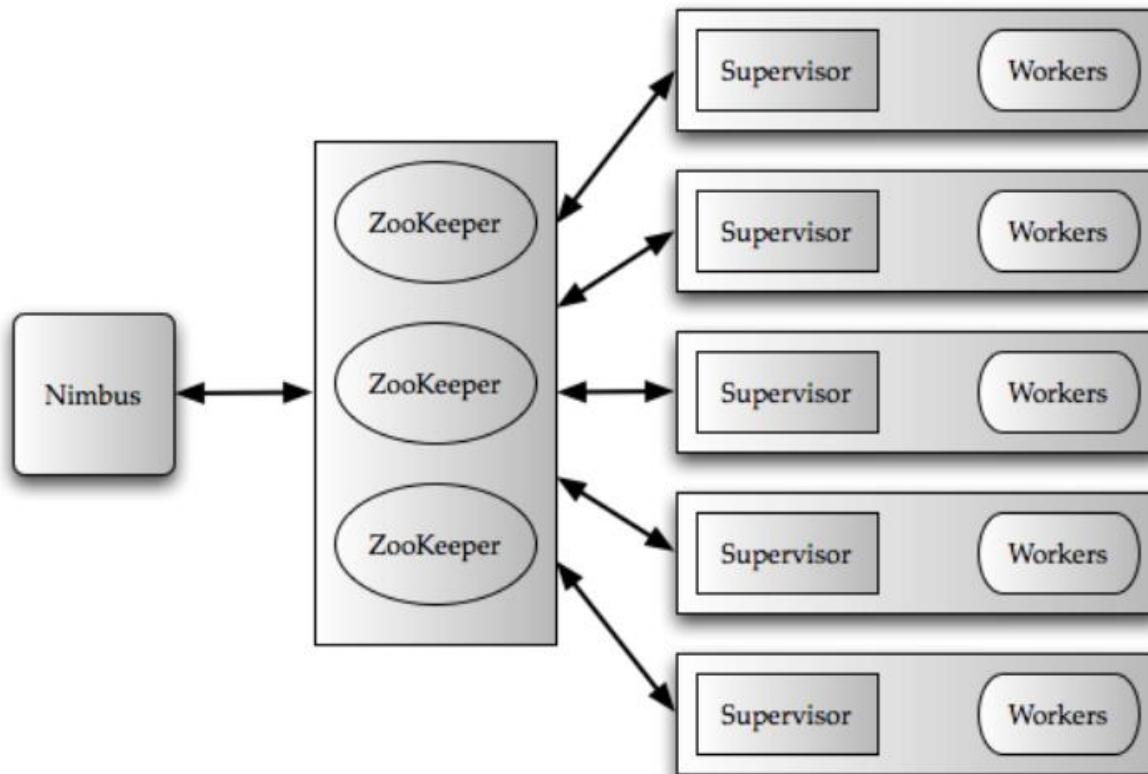
- Micro-batch: process tuples in small batches

- Spark Streaming



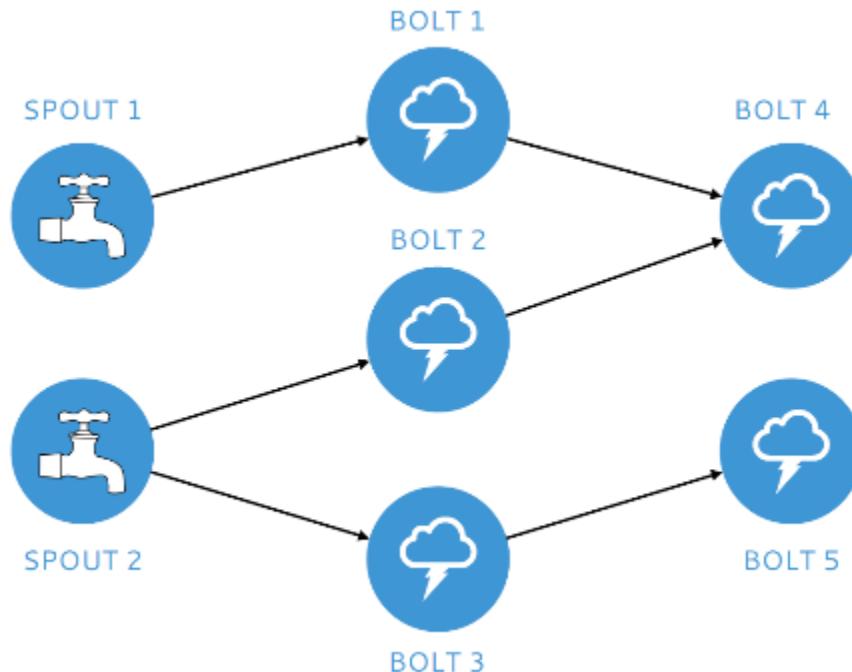
Apache Storm

- Storm is a classic **continuous dataflow** based distributed stream processing system developed by Twitter. It has many variations such as JStorm in Alibaba.



DAG model: topology

- Two kinds of components
 - Spout: read streams from sources and generates tuples
 - Bolt: execute a simple user defined function for tuples
- Each spout and bolt can generate multiple instances as multiple tasks and partition them in different nodes



Use case: word count on Storm

□ Topology Building

```
TopologyBuilder tb = new TopologyBuilder();
```

Specify the degree of parallelism

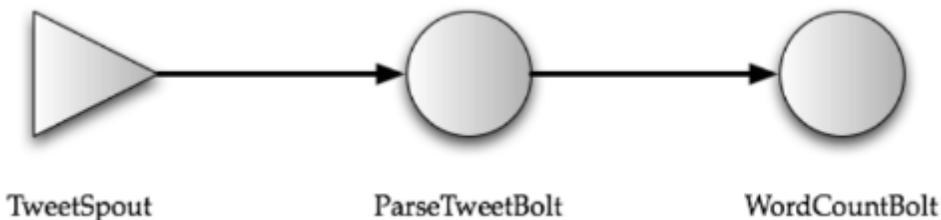
```
tb.setSpout("TweetSpout", kafkaSpout(), 4);
```

```
tb.setBolt("ParseTweetBolt", new ParseBolt(), 4).setNumTasks(8).shuffleGrouping("TweetSpout");
```

```
tb.setBolt("WordCountBolt", new CountBolt(), 4).fieldsGrouping("Split", new Fields("ParseTweetBolt"));
```

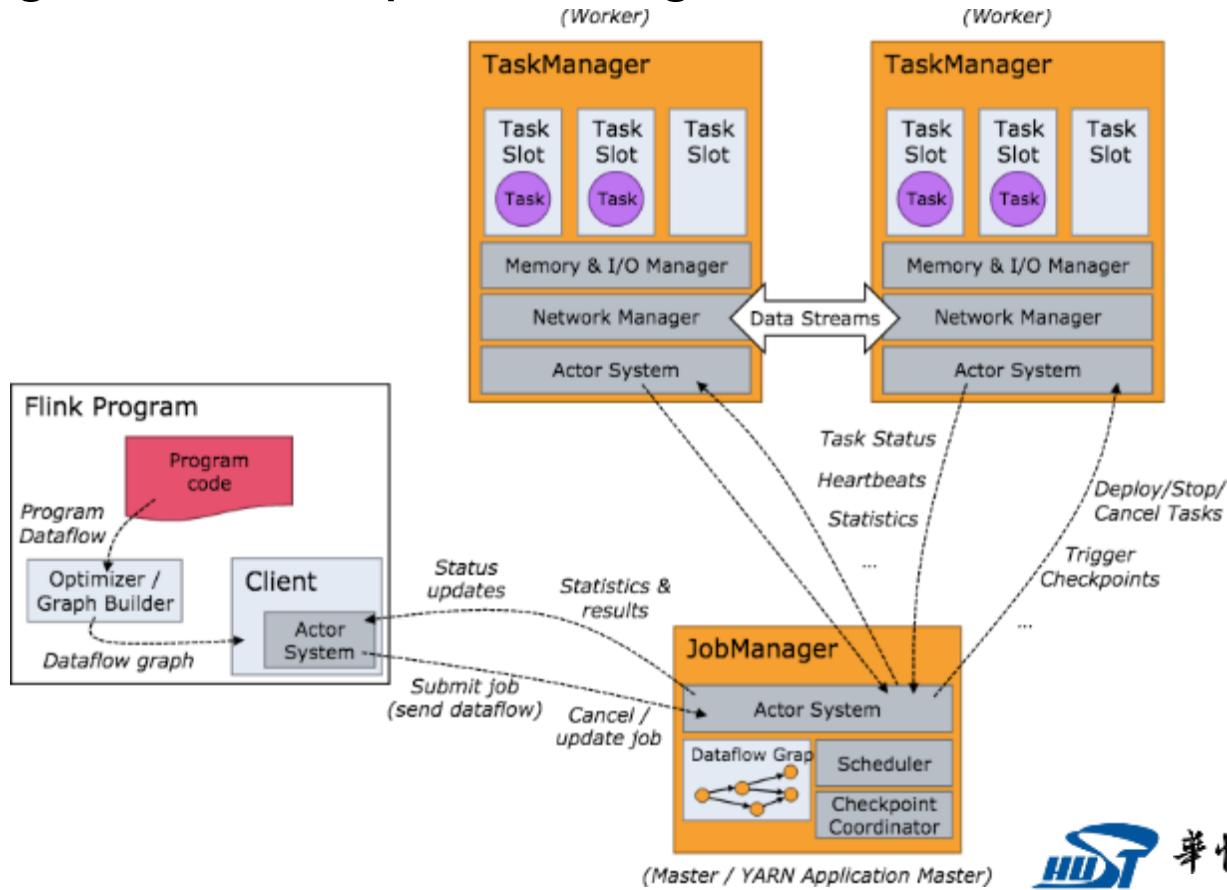
User defined functions

Specify the grouping strategies



Apache Flink

- Flink is one of the most popular **continuous dataflow** based distributed stream processing system, which can also support big data batch processing



DAG model: stream & transformation

- Flink's DAG model is almost like that in Storm, which contains one **source** operator, one **sink** operator and multiple **transformation** operators
 - Transformation can be user defined functions or many pre-defined operations provided by Flink API, e.g., flatmap, map, reduce, windows, ...



Use case: word count on Flink

```

DataStream<String> lines = env.addSource(
    new FlinkKafkaConsumer<>(...));
}

```

} *Source*

```

DataStream<Event> events = lines.map((line) -> parse(line));
}

```

} *Transformation*

```

DataStream<Statistics> stats = events
    .keyBy("id")
    .timeWindow(Time.seconds(10))
    .apply(new MyWindowAggregationFunction());
}

```

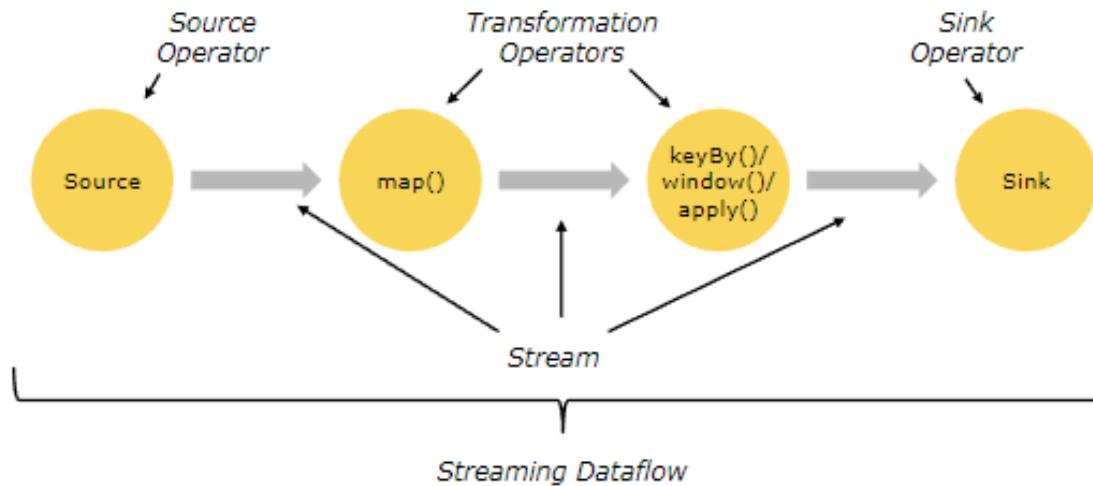
} *Transformation*

```

stats.addSink(new RollingSink(path));
}

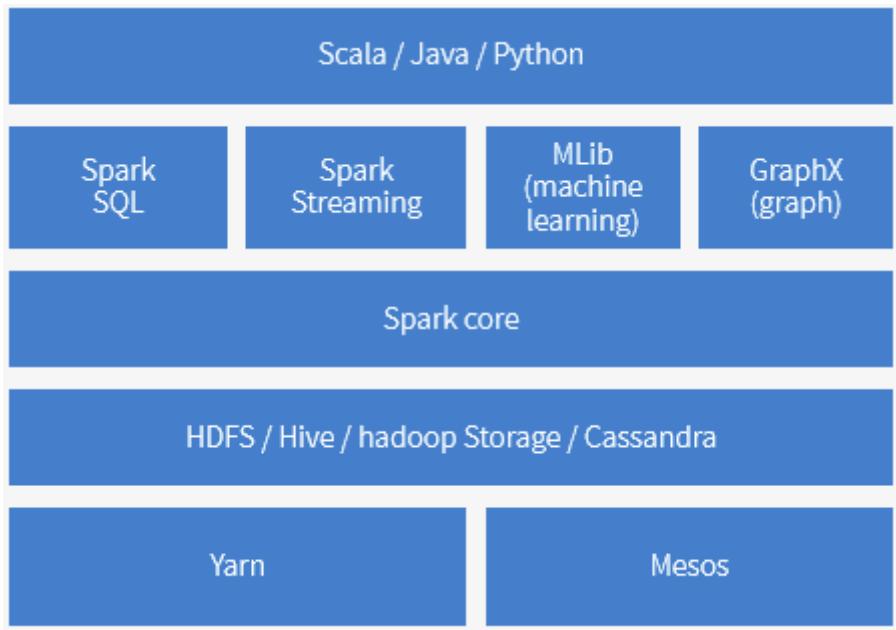
```

} *Sink*



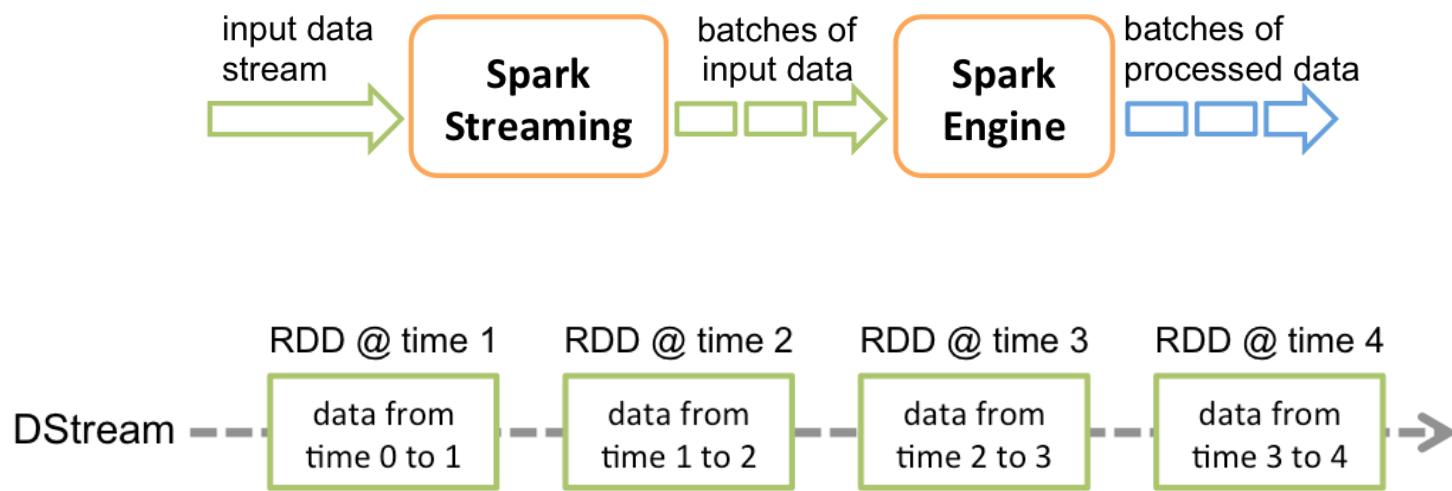
Spark Streaming

- Spark Streaming is an extension of the core Spark API, it is a classical **micro-batch** based distributed stream processing system



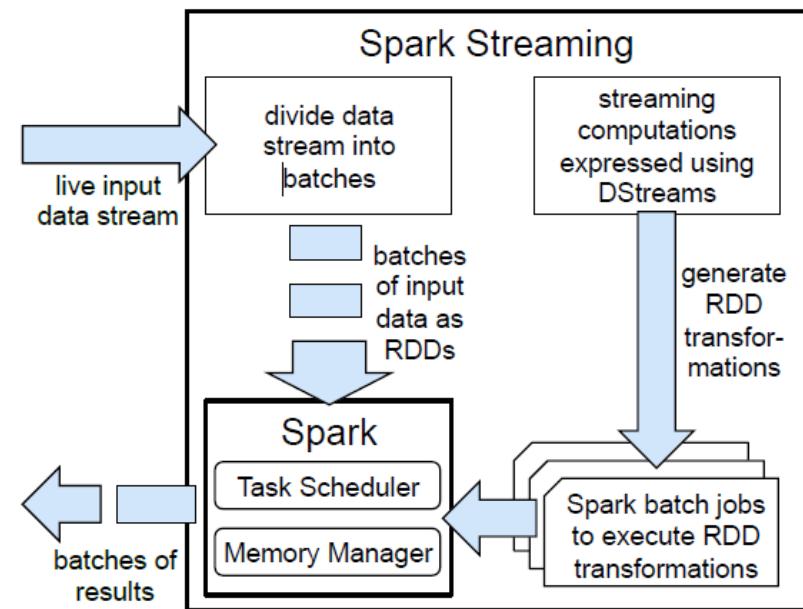
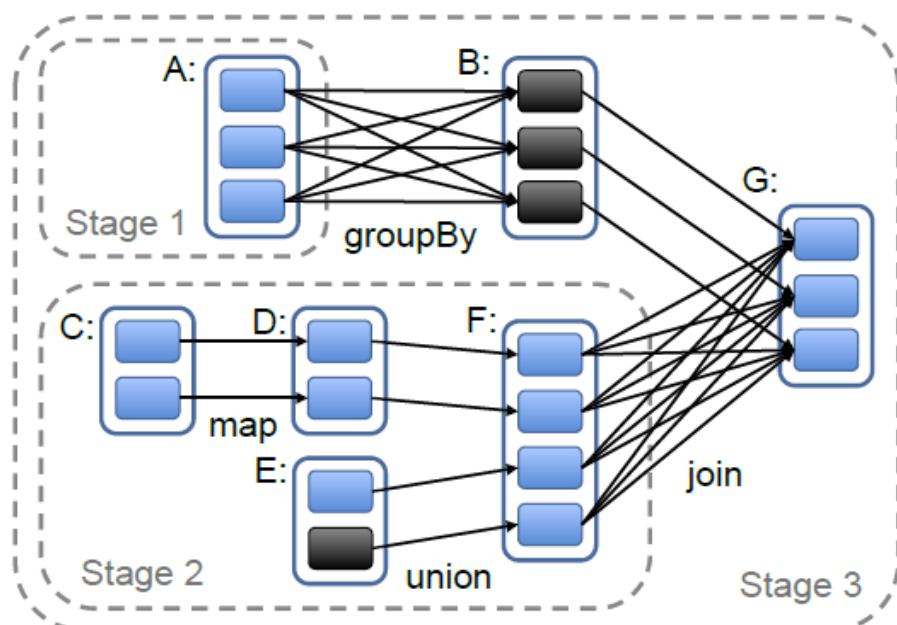
Discretized streams and RDDs

- The basic data abstraction of Spark Streaming is **discretized streams**
 - Tuples in a fixed interval (e.g., 1s) are partitioned into a small batch, and transformed to RDD by using Spark engine



DAG model: RDD transformations

- Spark Streaming executes RDD transformations for each small batch of tuples and generates batches of results



Word count on Spark Streaming

- ❑ Essentially the same as batch processing

```
JavaRDD<String> textFile = sc.textFile("hdfs://...");  
JavaPairRDD<String, Integer> counts = textFile  
    .flatMap(s -> Arrays.asList(s.split(" ")).iterator())  
    .mapToPair(word -> new Tuple2<>(word, 1))  
    .reduceByKey((a, b) -> a + b);  
counts.saveAsTextFile("hdfs://...");
```